

Java 101 for Capital One

Version 2.0.5 October 2014



Copyright © Descriptor Systems, 2015
All Rights Reserved.

This document may not be copied or reproduced in any form without prior written consent of Joel Barnum
of Descriptor Systems.

All trademarks belong to their respective companies.

You can contact Descriptor Systems at:

<http://www.descriptor.com>
sales@descriptor.com
319-362-3906
P.O. Box 461
Marion, IA 52302

Java 101 for Capital One

Chapter 1: Introduction to the Course

Java 101 for Capital One.....	1-1
Legal Information.....	1-2
Java 101 for Capital One.....	1-3
Introductions.....	1-4
Course Description.....	1-5
Course Objectives.....	1-6
Sample Agenda.....	1-7
Sample Agenda, cont'd.....	1-8

Chapter 2: Introduction to Java

Introduction to Java.....	2-1
What is Java?.....	2-2
Java History.....	2-3
Java Language Family Tree.....	2-4
Java Programming Platforms.....	2-5
Java Standard Edition Version History.....	2-6
Java Standard Edition Version History, cont'd.....	2-7
Portability.....	2-8
Java Development Process.....	2-9
The Java SDK versus the JRE.....	2-10
The Java Class Library.....	2-11
Introduction to JEE.....	2-12
Hello, World.....	2-13
Chapter Summary.....	2-14

Chapter 3: Data Types and Assignment

Data Types and Assignment.....	3-1
Java Program Structure.....	3-2
Java Comments.....	3-3
Java Statements.....	3-4
Variables.....	3-5

Rules for Identifiers.....	3-6
Reserved Words.....	3-7
Java Data Types.....	3-8
Integer Types.....	3-9
Floating Point Types.....	3-10
Character Type.....	3-11
Character Escape Codes.....	3-12
Boolean Type.....	3-13
The String Type.....	3-14
Assignment.....	3-15
Cast Operator.....	3-16
Scope of Variables.....	3-17
Scope of Variables, cont'd.....	3-18
The Stack and the Heap.....	3-19
Quick Practice.....	3-20
Chapter Summary.....	3-21

Chapter 4: Operators

Operators.....	4-1
Operators.....	4-2
Math Operators.....	4-3
Math Operators, cont'd.....	4-4
Other Math Operations.....	4-5
Quick Practice.....	4-6
Compound Assignment Operators.....	4-7
Increment and Decrement Operators.....	4-8
Equality Operators.....	4-9
Relational Operators.....	4-10
Integer Bitwise Operators.....	4-11
Boolean Logical Operators.....	4-12
Conditional Operators.....	4-13
Shift Operators.....	4-14
Conditional Operator.....	4-15
The Boolean ! Operator.....	4-16

String Concatenation.....	4-17
Other Operators.....	4-18
Operator Precedence.....	4-19
Chapter Summary.....	4-20

Chapter 5: Flow Control

Flow Control.....	5-1
Defining Blocks.....	5-2
Conditional Statements.....	5-3
The if Statement.....	5-4
The if else Statement.....	5-5
The if else if Statement.....	5-6
The if else if Statement, cont'd.....	5-7
Quick Practice.....	5-8
The while Loop.....	5-9
The while Loop, cont'd.....	5-10
The for Loop.....	5-11
The for Loop, cont'd.....	5-12
The for Loop, cont'd.....	5-13
Breaking Out of a Loop.....	5-14
The switch-case Statement.....	5-15
The switch-case Statement, cont'd.....	5-16
The switch-case Statement, cont'd.....	5-17
Quick Practice.....	5-18
Chapter Summary.....	5-19

Chapter 6: Introduction to Object Orientation

Introduction to Object Orientation.....	6-1
What is Object-Oriented Programming?.....	6-2
Object-Oriented Languages.....	6-3
Goals of Object Orientation.....	6-4
The Three Pillars.....	6-5
What is an Object?.....	6-6
Why are Objects Useful?.....	6-7
Objects Respond to Messages.....	6-8

Objects and Classes.....	6-9
A Class is Like a Cookie Cutter.....	6-10
Writing a Class.....	6-11
Sample Java Class.....	6-12
Instantiating Objects.....	6-13
Methods Operate on Objects.....	6-14
Methods Must Have a Receiver.....	6-15
Implicit Receivers.....	6-16
The Stack and the Heap.....	6-17
Comparing References.....	6-18
Comparing References, cont'd.....	6-19
Garbage Collection.....	6-20
Introduction to Inheritance.....	6-21
Introduction to Polymorphism.....	6-22
Synonym Alert!.....	6-23
Introduction to Packages.....	6-24
Encapsulation.....	6-25
Using Encapsulation.....	6-26
Java Encapsulation Summary.....	6-27
Get/Set Methods.....	6-28
Get/Set Methods, cont'd.....	6-29
Introduction to Interfaces.....	6-30
Quick Practice.....	6-31
Chapter Summary.....	6-32

Chapter 7: Methods

Methods.....	7-1
What is a Method?.....	7-2
Calling a Method.....	7-3
Method Receivers.....	7-4
Class Structure.....	7-5
Method Syntax.....	7-6
Instance Method Examples.....	7-7
Calling Instance Methods.....	7-8

Calling Methods, cont'd.....	7-9
Quick Practice.....	7-10
Quick Practice 2.....	7-11
Overloading Methods.....	7-12
Calling Overloading Methods.....	7-13
Constructor Methods.....	7-14
The new Operator.....	7-15
Overloaded Constructors.....	7-16
Calling Constructors.....	7-17
Quick Practice.....	7-18
The No-Argument Constructor.....	7-19
Calling One Constructor From Another.....	7-20
Field Initializers.....	7-21
Method Modifiers.....	7-22
Static Methods.....	7-23
Static Methods, cont'd.....	7-24
Implicit Receiver for Static Method.....	7-25
Static Fields.....	7-26
Static Fields, cont'd.....	7-27
Quick Practice.....	7-28
The final Keyword.....	7-29
Chapter Summary.....	7-30

Chapter 8: Exception Handling

Exception Handling.....	8-1
Traditional Error Handling.....	8-2
Traditional Error Handling Issues.....	8-3
Java Exception Handling.....	8-4
Advantages of Exception Handling.....	8-5
Uncaught Exceptions.....	8-6
Exceptions Are Objects.....	8-7
Multiple Catch Blocks.....	8-8
Handle the Most Specific Exception First.....	8-9
Passing Exceptions Back.....	8-10

The finally Block.....	8-11
Throwing Exceptions.....	8-12
Writing an Exception Class.....	8-13
Writing an Exception Class, cont'd.....	8-14
Quick Practice.....	8-15
Chapter Summary.....	8-16

Chapter 9: Arrays and Collections

Arrays and Collections.....	9-1
What is an Array?.....	9-2
Using Arrays.....	9-3
Array Bounds Checking.....	9-4
Arrays of Object References.....	9-5
Initializing Arrays.....	9-6
Quick Practice.....	9-7
The Enhanced for Loop.....	9-8
Introduction to Collections.....	9-9
Collections Can Automatically Grow.....	9-10
Basic Collection Types.....	9-11
Concrete Collection Classes.....	9-12
Data Structures.....	9-13
Using an ArrayList.....	9-14
Using an ArrayList, Pre Java 5.....	9-15
Using a LinkedList.....	9-16
Using a HashSet.....	9-17
Using a HashMap.....	9-18
Quick Practice.....	9-19
Chapter Summary.....	9-20

Chapter 10: Inheritance and Polymorphism

Inheritance and Polymorphism.....	10-1
What is Inheritance?.....	10-2
Why Use Inheritance?.....	10-3
Inheritance Vs Composed-Of Relationships.....	10-4
Java and Inheritance.....	10-5

Subclass Objects are a Superset.....	10-6
Accessing Superclass State and Behavior.....	10-7
Building Class Hierarchies.....	10-8
Quiz: What Kind of Relationship?.....	10-9
Quiz: What Kind of Relationship?.....	10-10
Quiz: What Kind of Relationship?.....	10-11
Overriding Behaviors.....	10-12
The Object Class.....	10-13
Calling Methods.....	10-14
Constructors and Superclasses.....	10-15
Constructors and Superclasses, cont'd.....	10-16
Polymorphism.....	10-17
Polymorphic Reference Assignment.....	10-18
Writing Polymorphic Algorithms.....	10-19
Using Abstract Classes.....	10-20
The final Keyword.....	10-21
Quick Practice.....	10-22
Review Questions.....	10-23
Chapter Summary.....	10-24

Chapter 11: Interfaces

Interfaces.....	11-1
What is an Interface?.....	11-2
Java Supports Single Inheritance.....	11-3
We Still Need Polymorphism.....	11-4
Solution: Interfaces.....	11-5
Defining an Interface.....	11-6
Implementing an Interface.....	11-7
Implementing Multiple Interfaces.....	11-8
'Tagging' Interfaces.....	11-9
Interface Inheritance.....	11-10
Implementing a Sub-Interface.....	11-11
Convenience Adapter Classes.....	11-12
Interfaces versus Abstract Classes.....	11-13

Interface Reference Types.....	11-14
Quick Practice.....	11-15
Interfaces as Contracts.....	11-16
Examples of Using Interfaces.....	11-17
Example: Using the Comparable Interface.....	11-18
Using the Comparable Interface, cont'd.....	11-19
Using the Comparable Interface, cont'd.....	11-20
Using the Comparable Interface, cont'd.....	11-21
Using the Comparable Interface, cont'd.....	11-22
Example: Using the ContentHandler Interface.....	11-23
Example: Using ContentHandler, cont'd.....	11-24
Example: Using ContentHandler, cont'd.....	11-25
Review Questions.....	11-26
Chapter Summary.....	11-27

Java 101 for Capital One

Version 2.0.5 October 2014



1 - 1

Legal Information

Copyright (C) Descriptor Systems 2001, 2014

All Rights Reserved

All trademarks are owned by their respective companies

You can contact Descriptor Systems at:

<http://www.descriptor.com>

jbarnum@descriptor.com

1 - 2

Java 101 for Capital One

- Introductions
- Objectives

1 - 3

Introductions

- Please introduce yourself:
 - Name:
 - Java background (if any):
 - Other programming languages you know:
 - Your goal for the course:

1 - 4

Course Description

Description: This course introduces the student to writing object-oriented programs in Java.

Prerequisites: Prior study in object-orientation and UML is helpful.

Audience: Business analysts, developers, managers and other people interested in learning how to program Java.

Course Objectives

After taking this course, you will be able to:

- Write Java classes using object-oriented techniques such as encapsulation, inheritance and polymorphism
- Write Java programs using basic syntax elements for looping and flow of control
- Write Java programs that define and manipulate standard Java data types
- Write Java programs that create and manipulate Java arrays and collections
- Write Java classes with constructors, overloaded methods and static members
- Write Java programs that catch and throw exceptions

Sample Agenda

Day 1

Introduction to the Course

Introduction to Java

Data Types and Assignment

Day 2

Operators

Flow Control

Object-Oriented Programming Fundamentals

1 - 7

Sample Agenda, cont'd

Day 3

Methods

Exception Handling

Day 4

Arrays and Collections

Inheritance and Polymorphism

Interfaces

1 - 8

Introduction to Java

- What is Java?
- Hello, World

2 - 1

What is Java?

- Java is an object-oriented programming language created by Sun Microsystems, now owned by Oracle
- But Java is more than just a language -- it defines a portable platform for developing software that can execute in many environments



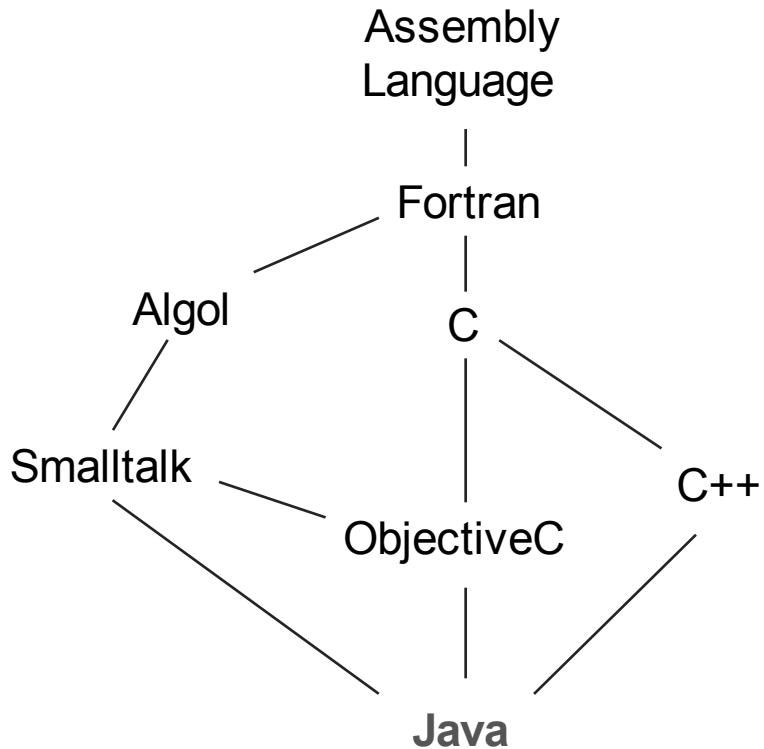
2 - 2

Java History

- Java was created at Sun in the early 1990s by James Gosling, who was initially attempting to create a processor-independent language for TV set-top boxes
- Java's original name was "Oak" after the tree outside of Gosling's office window
- First working application was a Web browser named WebRunner, later renamed to HotJava

2 – 3

Java Language Family Tree



2 - 4

Java Programming Platforms

- *Java Standard Edition (JSE)* provides the core Java functionality for client and server-side programming
- *Java Enterprise Edition (JEE)* extends J2SE so developers can create server-side applications (e.g. Web applications)
- *Java Micro Edition (JME)* is a specialized Java for small devices (e.g. Cellphones and personal digital assistants)

Java Standard Edition Version History

Version 1.0

Basic class library, AWT GUI toolkit

Version 1.1

Added new I/O libraries, event-model enhanced internationalization support and inner classes

Version 1.2

New security model, Java Foundation Classes including the Swing GUI toolkit

Version 1.3

Mostly performance enhancements and fixes

Version 1.4

XML support, logging

Java Standard Edition 5

Generics, annotations, formatted I/O

Java Standard Edition 6

Client-side Web services, built-in database

2 - 6

Actually, the term "Standard Edition" didn't come about until Java version 1.2, which Sun's marketing department dubbed "Java 2".

Java Standard Edition Version History, cont'd

Version 7.0

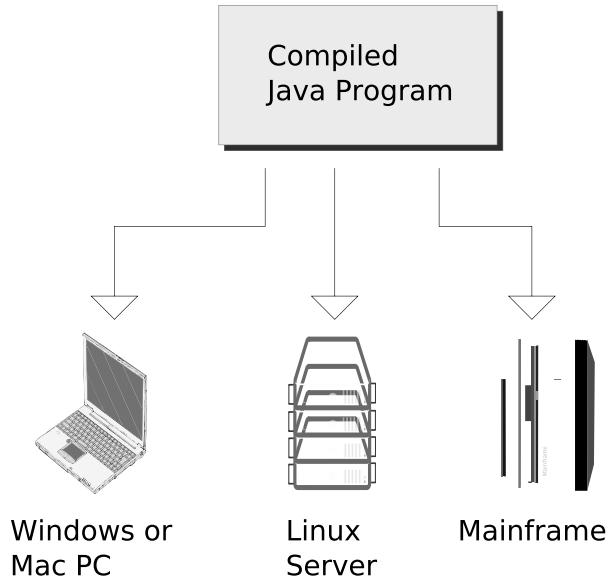
try-with-resources, strings in switch-case

Version 8.0

Functional programming, parallel programming

Portability

- Java's motto is *Write Once, Run Anywhere* -- compiled Java programs should run unchanged on different operating systems and computing platforms



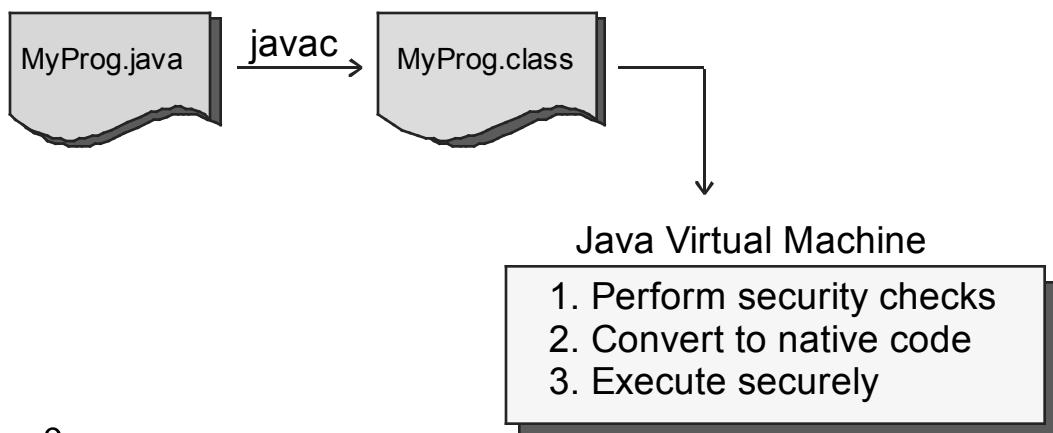
2 - 8

Portability is one of the primary advantages of Java over other languages -- note that we are talking about executable portability, not source-code portability.

Of course, this picture depends on the Java platform itself being ported to each of the environments.

Java Development Process

- To develop a Java program, you create a source file, then use the Java compiler to create a file containing *byte codes*
- You then pass the byte codes to a *Java Virtual Machine* that converts the byte codes to native code and executes



2 – 9

Java programs are compiled, but not into native code. Instead, the compiler produces an intermediary form known as "byte codes". One of the JVM's jobs is to translate the byte codes into native codes on a particular operating system and CPU. So for Java to run on any operating system or CPU, someone must port the JVM to it.

Sun has published the specification for the JVM -- anyone can implement it. There currently are several JVMs available from Sun, IBM and other vendors.

The Java SDK versus the JRE

- To run Java programs, you need the Java Runtime Environment, which includes the JVM
- To develop Java programs, you need a Java 2 Software Development Kit
- The Oracle Java SDK includes a JRE

2 – 10

The JRE also includes accessory files such as fonts.

The Java Class Library

- Java ships with a portable set of pre-written code known as the Java class library

User-interfaces (AWT and Swing)
Network I/O (e.g. sockets, HTTP)
File I/O
Threads
Security
Internationalization
Collection data structures
XML processing

....

2 – 11

This is another strength of Java -- unlike languages like C++, which have rudimentary (at best) library support, the JRE contains a very large (and sometimes intimidating) set of functions that you can use in your Java programs.

Introduction to JEE

- Java Enterprise Edition is a set of specifications for enterprise computing
- Most JEE technologies focus on server-side development

Servlets

JavaServer Pages

Enterprise JavaBeans

Java Naming and Directory Interface

JDBC

Java Message Service

...

2 – 12

JEE previously was known as Java 2 Enterprise Edition (J2EE).

Hello, World

Test.java

```
1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          System.out.println ("Hello, world");
6      }
7  }
```

```
c:\> javac Test.java
c:\> dir
03/01/2005  08:16 AM           414 Test.class
03/01/2005  08:16 AM           115 Test.java
c:\> java Test
Hello, world
```

2 - 13

Chapter Summary

In this chapter, you learned:

- What Java is all about

Data Types and Assignment

- Defining Variables
- Primitive Data Types
- Assignment

3 - 1

Java Program Structure

- The **class** is the basic building block of a Java program
- Classes define state (variables) and behaviors (methods)

```
1  public class MyClass
2  {
3      private int x;    // class-scoped variable
4
5      public static void main(String[] args)
6      {
7          int z = 12;        // local variable
8
9          System.out.println("z: " + z);
10     }
11 }
```

3 – 2

In later chapters, you will learn more about classes, but we show a sample class here so you can see the overall structure of a Java program.

Classes define variables to hold data and define methods (a.k.a. functions or procedures) in which you write Java code.

The "main()" method shown here is special since it acts as the entry point of the class.

Java Comments

- Java supports three kinds of comments:
 - Single line comments
 - Block comments
 - JavaDoc comments

```
1  /**
2   * @author Paul Westerberg
3   * @version 1.0
4   */
5  public class Test
6  {
7      // a single-line comment
8      /* a block comment
9       * of multiple lines
10     */
11 }
```

3 – 3

Single-line comments comment-out until the end of the line.

Block comments comment-out until the end of the comment block.

JavaDoc comments are special comments used to create automatically generated documentation for a Java class. The J2SDK includes a javadoc utility that scans source files for JavaDoc comments and generates HTML documentation. Note that the source JavaDoc comments can contain special tags that have meaning to the javadoc processor. For more information on JavaDoc, see:

<http://java.sun.com/j2se/javadoc/writingdoccomments>

Java Statements

- In general, Java programs consist of **statements**, each of which is terminated with a semicolon
- Statements must reside within a **block** defined with curly braces

```
1  public class MyClass
2  {
3      private int x;    // class-scoped variable
4
5      public static void main(String[] args)
6      {
7          int z = 12;        // local variable
8
9          System.out.println("z: " + z);
10     }
11 }
```

3 - 4

Lines 3, 7 and 9 are examples of statements. Note that statements can span multiple lines -- it's the semicolon that terminates the statement.

This program defines two blocks: one that goes from line 2 until line 11, the other is a nested block from lines 6 to 10. You can nest blocks indefinitely as required.

Variables

- **Variables** are locations in memory that hold a value
- In Java, you define a variable with a statement that specifies both a type and a name (a.k.a an identifier)
Basic syntax:

Type Name [=initial value];

```
1 int count = 12;  
2 float interestRate;  
3 String SQL="SELECT * FROM employee";
```

3 – 5

Any programming language needs variables to hold data for processing. In Java, the basic syntax to define a variable requires you to specify the variable's type, followed by its name, followed by an optional initialization value.

The three examples shown here demonstrate defining variables that contain an initialized integer, an un-initialized floating-point number and a String of characters.

Rules for Identifiers

- Any length combination of letters, digits and special characters: ('_ ', '\$') but must begin with an letter or special character
- Java is case sensitive!

Valid

count
interestRate
purchase_Number1

Invalid

1account
order-number
delivery date

3 – 6

The rules for identifiers are straightforward. You can (and should) use names that are descriptive and you can use underscores or uppercase characters to separate words within identifiers.

Reserved Words

abstract	final	outer	try
boolean	finally	package	var
break	float	private	void
byte	for	protected	volatile
byvalue	future	public	
case	generic	rest	
cast	goto	return	
catch	if	short	
class	implements	static	
clear	import	strictfp	
const	inner	super	
continue	instanceof	switch	
default	int	synchronized	
do	interface	this	
double	long	throw	
else	native	throws	
extends	new	transient	
false	null	true	

3 – 7

This page lists the reserved identifiers that you must not use when you assign names to your variables and other identifiers. Some of these keywords are reserved, but not implemented (e.g. goto).

Java Data Types

- Java is a strongly typed language -- when you perform an assignment, the compiler flags an error if the types are not compatible
- There are two basic kinds of types in Java: **primitive** types and **reference** types

Primitive types:

byte, short, int, long
float, double
char
boolean

Reference types:

java.lang.String
java.util.Date
mypackage.BankAccount
...

3 – 8

Primitive types are basic types such as characters, integers and floating-point numbers. Reference types are more complex data, such as Dates, Strings, Connections and so forth. As we will see later, variables of primitive types hold values, while variables of reference types hold a reference to the actual storage for the data.

All reference types are defined in a "package", e.g. java.util. Packages let you organize the types into categories. You can omit the package name for types defined in the "java.lang" package.

Integer Types

- Integers are primitive types that contain whole numbers
- Literal constants by default take on the *int* type

Type	Size (bits)	Size (bytes)	Values	Example
byte	8 bits	1 byte	-128 to 127	byte b = 12;
short	16 bits	2 bytes	-32768 to 32767	short sh = 1000;
int	32 bits	4 bytes	-2**31 to 2**31-1	int i2 = 45900;
long	64 bits	8 bytes	-2**63 to 2**63-1	long l = 1000000;

3 – 9

Integers are the workhorse type for most Java programs. Note that there are four sizes of integers that let you trade off storage space and capacity.

Unlike in languages like C, all Java types, including the integer types have a fixed size and range on all platforms. That makes Java more portable.

Floating Point Types

- Floating point values are primitives that contain both whole and fractional parts
- By default, floating point literal constants are of type double

Type	Size (bits)	Size (bytes)	Values	Example
float	32 bits	4 bytes	1.4e-45 to 3.4e+38	float f = 12.23F;
double	64 bits	8 bytes	4.9e-324 to 1.7e+308	double d = 26.77e3;

3 – 10

Java floating point types adhere to the IEEE standard for floating point numbers.

Floating-point literal constants with no specifiers (e.g. 12.34) take on the double type.

The `java.lang.Math` package contains useful methods and fields to work with floating-point values (e.g. `Math.PI`).

Character Type

- The `char` type holds a single 16-bit Unicode character
- To form a character literal constant, use single quotes

```
1     char ch1 = 'A';
2     char ch2 = "B"; // won't compile
```

3 – 11

Unicode is an encoding scheme that maps each character from most (if not all) national languages into a unique number. Unicode uses at least 16 bits to perform this mapping and thus can map far more characters than the old ASCII and EBCDIC encoding schemes. Since Java uses Unicode internally, Java is well positioned to let you write internationalized (I18N) programs.

For more information on Unicode, see www.unicode.org.

Character Escape Codes

- Java defines special escape codes for common non-printable characters
- You can initialize a character to any Unicode value using an integer constant

\b	Backspace
\t	Horizontal tab
\n	Linefeed
\f	Form feed
\r	Carriage return
\"	Double quote
'	Single quote
\\	Backslash
\uxxxx	Hexadecimal Unicode code point
\xxx	Octal Unicode code point

3 – 12

You can specify any Unicode character using the character's code point and the '\uxxxx' notation, but common characters have shortcuts (e.g. \n). Examples:

```
char ch1 = '\n'; char ch2 = '\u00DC'; // German U umlaut
```

Boolean Type

- The boolean type allows only the literal values true and false
- The most common use of booleans is in flow control statements (e.g. the "if" statement)

```
boolean finished = false;  
boolean ready = true;
```

3 – 13

Java programmers often use boolean variables in loops or conditional statements (e.g. "if" statements). The keywords "true" and "false" are considered literal boolean constants.

The String Type

- Strings are reference types that hold sequences of characters
- To define a literal String, use double quotes
- You can concatenate Strings using the + operator

```
String s1 = "If Only You Were Lonely";
int index = s1.indexOf ( "Only" );
boolean same = s1.equals ( "The Ledge" );

String s2 = s1 + " Not";
```

3 – 14

Unlike the previously covered types, Strings are object types, not primitives. That means that Strings are "smarter" than primitives -- for example, Strings support searching, tests for equality, extraction of sub-strings and so forth.

As we will cover later in the course, String is actually an object-oriented class that's part of the standard Java class library.

When you initialize a variable for an object type, the variable holds a reference to the object, not the object's value.

Assignment

- The = operator assigns the right-side value to the left-side
- The right-side data type must be compatible with the left side, or else the compile flags an error

```
1  int i1 = 123456;
2  short s1 = 12;
3  float f1 = 12.34F;
4  double d1 = 1.4e87;
5  String str1 = "ABCD";
6
7  int i2 = s1;    // OK
8  short s2 = i1; // compile error
9
10 double d2 = f1; // OK
11 float f2 = d1; // error
12
13 i2 = str1;     // compile error
```

3 – 15

For an assignment to be avoid compile errors, the left-side type must be at least as "large" as the right-side type.

Note that it's never OK to assign a reference type (e.g. String) to a primitive variable.

Cast Operator

- To force an assignment, you can use the cast operator on the right side of an expression
- You should avoid casting if you can!

```
1 int i1 = 123456;
2 short s1 = 12;
3 float f1 = 12.34F;
4 double d1 = 1.4e87;
5
6 short s2 = (short)i1;    // compiles OK, truncates
7
8 float f2 = (float)d1;   // compiles OK,
9                      // result is Infinity
10                     // (overflow)
```

3 – 16

The cast operator converts from one type to another. To form a cast, you enclose the target type within parenthesis.

Casting does help in some situations to avoid compile errors, but may result in odd runtime behavior.

It is possible to cast from one reference type to another, but not to/from a reference type and a primitive.

Scope of Variables

- Variables must be defined within a *block* defined by curly braces { }
- Variables defined within a block are local to that block

```
{  
    int i = 12;  
    j = i;      // compile error  
  
    {  
        int j = 13;  
        int k = i;  // OK  
    }  
}
```

3 – 17

You can define a block by using a pair of curly braces -- any variables defined in a block are local to that block (and any sub-block). Code within a sub-block can access variables defined in any outer block.

The code shown here is a contrived example to show the syntax, but it's quite common to nest blocks within other blocks, especially using looping or conditional statements.

Scope of Variables, cont'd

- Variables defined at *class scope* are global to all methods (functions) in the class

```
1  public class Test
2  {
3      private int i1 = 15; // field
4
5      public void method1()
6      {
7          int i2 = 14; // local variable
8          i2 = i1 + 5; // OK
9      }
10     public void method2()
11     {
12         int i3 = i2; // compile error
13         int i4 = i1 + 5; // OK
14     }
15 }
```

3 – 18

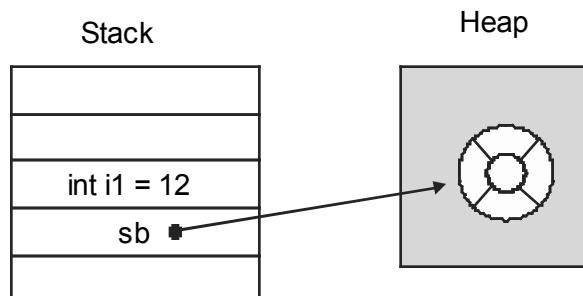
The code on this page shows a Java class definition -- we will cover the syntax in more detail later, but please note that a class defines a block using curly braces. In lines 5 to 9, we nest a method (a.k.a. function) definition which also defines a block. A class can contain many methods, each with its own block -- in lines 10–14 we define another method.

In line 3, we define a variable at "class scope" -- such variables are global to all methods within the class and are referred to as "fields". In contrast, line 7 defines a local variable, "i2", which is only available to the method in which it's defined -- that's why line 12 generates a compile error.

The Stack and the Heap

- The Java Virtual Machine defines two basic areas for variables: the *stack* and the *heap*
- Local variables are allocated on the stack, while dynamically allocated storage comes from the heap

```
1  public class Test
2  {
3      public void method1()
4      {
5          int i1 = 12;
6          StringBuffer sb = new StringBuffer();
7      }
8  }
```



3 – 19

The "method1" method defines two local variables: i1 and sb. When the method is called, these two variables will be allocated on the JVM's stack and will be de-allocated when the method finishes.

StringBuffer, like the String type covered earlier, is a reference type that's part of the Java standard class library. Lines 6 and 7 allocate a StringBuffer object using Java's "new" operator, which allocates storage for the StringBuffer object from the JVM's "heap", which is the free storage in the JVM.

Note that the sb variable, like all reference variables, does not contain the actual StringBuffer object, but instead references, or points to the actual object on the heap.

Quick Practice

- Define an integer variable and initialize it to your favorite whole number:
- Define a floating point number variable and initialize it to your favorite floating-point number:
- Define a String variable and initialize it to your name:

3 – 20

Chapter Summary

In this chapter, you learned:

- The rules for defining and initializing variables
- The Java primitive data types
- About the stack and the heap

Operators

- Operator Types
- Assignment

4 - 1

Operators

- Operators let you perform math, comparisons and so forth
- Java provides 37 operators, which we can divide into categories:
 - Math
 - Equality and Relational
 - Assignment
 - Boolean Logical
 - Integer Bitwise
 - Shift

4 – 2

Math Operators

- Java defines math operators that work on numeric types (int, float and so forth)
- The compiler will promote operands to the larger type if necessary
- Integer division by zero causes an `ArithmaticException` to be thrown

+	Addition
-	Subtraction
*	Multiplication
/	Division
%	Remainder

4 – 3

These operators should be familiar to anyone who has taken an algebra course! The only one that might not be familiar is the remainder operator, which returns the remainder of integer division. For example, $7 \% 3 = 1$.

You should also note that just like in algebra, you can use parenthesis to force an math expression to evaluate in the order you desire.

Math Operators, cont'd

```
1  int i1 = 12;
2  double d1 = 12.34;
3
4  System.out.println ( i1 + d1 );
5
6  int i2 = 0;
7  try
8  {
9      System.out.println ( i1 / i2 );
10 }
11 catch ( ArithmeticException exc )
12 {
13     System.out.println ( "Math error! "+exc );
14 }
```

4 – 4

This code fragment demonstrates promotion of an integer to a double and also how to handle an `ArithmetiException`. We will cover exception handling in more detail later in the course.

Other Math Operations

- The `java.lang.Math` class defines additional math operations beyond the standard math operators (e.g. trig functions)
- These **methods** are not intrinsic Java syntax, but are part of the Java class library

```
1 double d = Math.sin ( Math.PI / 2 );
2 System.out.println ( d );
3
4 double radius = 5.3;
5 double area = Math.pow ( radius, 2 ) * Math.PI;
6 System.out.println ( area );
```

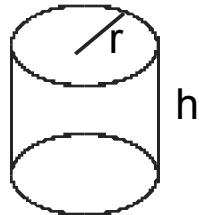
4 – 5

Just for completeness, we show here some of mathematical functions from the `Math` class. These are not actually part of the Java syntax -- instead, they are defined as methods in that class. Please read the `Math` documentation for a complete listing of the methods in this class.

Quick Practice

- The equation for the surface area of a cylinder is:

$$2\pi r^2 + 2\pi r h$$



- Complete this program fragment to calculate a cylinder's surface area:

```
1 double r = 12.0;  
2 double h = 3.0;  
3  
4 double surfaceArea =  
5  
6 System.out.println(surfaceArea);
```

4 - 6

Compound Assignment Operators

- Java provides several compound assignment operators

```
*= /= %= += -= <<= >>= >>>= &= ^= |=
```

```
1 int i1 = 12;
2
3 i1 += 45; // same as i1 = i1 + 45;
4 System.out.println ( i1 );
```

4 – 7

Compound assignment operators are shorthand syntax for performing an operation and an assignment. The example demonstrates using a compound assignment to perform an addition and assignment. With the exception of `+=` and Strings, all of the compound assignment operators work only on primitive types.

We will cover the other operators later in this chapter.

Increment and Decrement Operators

- These shortcut operators come in two forms:
prefix and **postfix**
- Postfix operators work after the expression is evaluated, prefix before

++ --

```
1 int i1 = 12;
2 i1++; // same as i1 = i1 + 1;
3 System.out.println ( i1 );
4
5 int i2 = --i1 * 3;
6 System.out.println ( i2 );
```

4 – 8

These operators work only on primitives and change the value of a variable. The interesting thing about them is that you can use them in expressions -- they are especially useful in looping constructs.

The example shows using the postfix ++ and prefix --.

Equality Operators

- Equality operators return a boolean result and are most often used in "if" statements or loops
- Equality tests of primitives test values, while equality tests of reference types tests the references

== !=

```
1 int i1 = 12;
2 int i2 = 14;
3
4 boolean b = i1 == i2;
5 System.out.println ( b );
6
7 if ( i1 != i2 )
8     System.out.println ( "Different!" );
```

4 – 9

These operators work on: primitive numeric types, booleans, reference types or the literal value "null". The operators always return a boolean type.

Relational Operators

- The numeric comparison operators work on primitive numeric types and return a boolean result
- The **instanceof** operator works on reference types and returns true only if the right-side is of the same type or subtype

```
< <= > >= instanceof  
1 int i1 = 12;  
2 int i2 = 14;  
3  
4 boolean b = i1 >= i2;  
5 System.out.println ( b );  
6  
7 if ( i1 < i2 )  
8     System.out.println ( "Smaller!" );
```

4 – 10

The relational operators let you compare numeric values or reference types and always return a boolean result. The numeric relational operators are especially handy for "if" statements and looping constructs.

The "instanceof" operator lets you check to see if a reference type is compatible with a specified type. It's quite useful when you store objects of different types in a collection -- when you iterate through the collection, you can use instanceof to determine the type.

Integer Bitwise Operators

- You can use integer bitwise operators to test or manipulate bit patterns within integer values (byte, short, int, long)

& ^ |

```
1 int i1 = 18;
2 int i2 = i1 & 8;    // 10010 and 01000
3 int i3 = i1 | 8;    // 10010 or 01000
4 int i4 = i1 ^ 16;   // 10010 xor 10000
5 System.out.println( i2 ); // 0
6 System.out.println( i3 ); // 26
7 System.out.println( i4 ); // 2
```

4 – 11

When you use these operators on integral numeric primitive types, they perform bit-level manipulation. They are not valid for characters, floating-point types or reference types. They do work on booleans as discussed on the next page.

Boolean Logical Operators

- You can use logical operators on booleans to build more complex testing expressions

& ^ |

```
1  boolean atHome = true;
2  boolean eatingDinner = false;
3
4  boolean result = atHome & eatingDinner;
5  System.out.println ( result );
6
7  result = atHome | eatingDinner;
8  System.out.println ( result );
9
10 result = atHome ^ eatingDinner;
11 System.out.println ( result );
```

4 – 12

These operators are most often used in "if" statements or looping constructs. For an "and" expression to be true, both operands must be true. For "or" to be true, either operand can be true. For exclusive or (^) to evaluate to true, the operands must be different.

Note that there are similar logical operators for and (&&) and or (||) shown on the next page.

Conditional Operators

- The **conditional operators** work similarly to the boolean logical operators, but they can "short-circuit"

&& ||

```
1 int i1 = 12;
2 int i2 = 13;
3
4 boolean result = (i1 < i2) || (++i1 == i2 );
5 System.out.println ( result );
6 System.out.println ( i1 );
7
8 result = (i1 > i2) && (++i1 == i2 );
9 System.out.println ( result );
10 System.out.println ( i1 );
```

4 – 13

We call these operators short-circuit operators because Java will not bother evaluating the second part of the expression if it doesn't need to. In the example, in line 4, the second expression containing the ++ operator will not execute, since an "or" expression is true if either operand is true and the first expression is true. In other words, Java can stop evaluating an "or" expression as soon as it finds a true value. Same for line 8, since Java can stop evaluating "and" as soon as it finds a false value.

While this might seem tricky, it can be useful in expressions like:

```
boolean result = ( myString != null ) && ( myString.length() > 5 );
```

Shift Operators

- You can use the shift operators to move bits within integer values (byte, short, int, long)
- You can use shift operators to multiply and divide by powers of two

<< >> >>>

```
1 int i1 = 18; // 10010
2 System.out.println ( i1 << 1 ); // 18 * 2 = 36
3 System.out.println ( i1 >> 2 ); // 18 % 4 = 4
4 int i2 = -18;
5 System.out.println ( i2 << 1 ); // -18 * 2 = -36
6 System.out.println ( i2 >> 2 ); // -18 % 4 = -4
7 System.out.println ( i2 >>> 2 ); // 1073741819
```

4 – 14

These operators perform a low-level operation on numeric integral types. In some cases performing multiplication and division in this fashion may be faster than using the multiply and divide operators.

The >>> operator performs a right-shift without preserving the integer's sign. For more on how this works, see any good computer programming text under the topic of "twos-complement" arithmetic.

Conditional Operator

- The conditional operator is a shortcut for an if-else statement

? :

```
1 int i1 = 12;
2 int i2 = 14;
3
4 int min = ( i1 < i2 ) ? i1 : i2;
5 System.out.println ( min );
```

4 – 15

This ternary (three operand) operator is a syntax that Java inherited from the C language. Java evaluates the first operand -- if it evaluates to true, the entire expression evaluates to the second operand. If false, the entire expression evaluates to the third operand.

The Boolean ! Operator

- The **not** operator flips the value of a boolean and is often used in "if" statements or loops

!

```
1  boolean atHome = true;
2  boolean eatingDinner = false;
3
4  boolean result = !(atHome && eatingDinner);
5  System.out.println ( result );
```

4 – 16

This unary (single-operand) operator has a high precedence, and thus is often used with parenthesis.

String Concatenation

- You can use the + and += operators to concatenate Strings

```
1  String s1 = "Hello";
2  String s2 = s1 + " world";
3
4  s1 += " Dolly";
```

4 – 17

The String class is the only non-primitive type that supports this operator. Under the covers, when you concatenate Strings, Java actually creates a new String with the additional characters and discards the original String (makes it available for garbage collection).

Other Operators

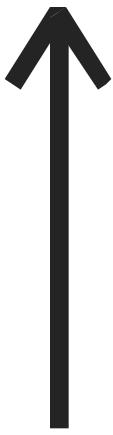
- . Object member access
- [] Array element access
- () Method call
- new Creates a new object or array from the heap
- () Cast operator
- "Flips" the sign of numeric values

4 – 18

For completeness, here we show the operators not covered explicitly in this chapter. We will see most of them later in the course.

Operator Precedence

Higher



Lower

postfix	[] .., (params), expr++, expr?
Unary	++expr, !expr, +expr, -expr!
Multiplicative	* / %
Additive	+ -
Relational	<, >, <=, >=
Equality	==, !=
Logical AND	&&
Logical OR	
Ternary	?
Assignment	=, +=, -=, *=, /=, %=

4 – 19

To force an order of precedence, use parenthesis.

Chapter Summary

In this chapter, you learned:

- About all of the Java operators
- How to create expressions that you can use in "if" statements and loops

Flow Control

- Conditional Statements
- Looping

5 - 1

Defining Blocks

- In Java, all executable code must be within a block defined by a pair of curly braces
- You can nest blocks within other blocks

```
1  public class Test
2  {
3      public void myMethod(int x)
4      {
5          if ( x == 12 )
6          {
7              // do something interesting
8          }
9      }
10 }
```

5 – 2

Conditional Statements

- Like all complete programming languages, Java provides syntax for making decisions:
 - if
 - Conditional (ternary) operator

5 – 3

Conditional statements let programs make decisions, generally by evaluating an expression that evaluates to a boolean. The basic conditional statement in Java is the "if" statement.

The if Statement

- The *if* statement examines an arbitrarily complex boolean expression

```
if ( <expr> )
{
    // statements
}
```

```
if ( <expr> )
    // one statement
```

5 – 4

The "if" statement evaluates the provided boolean expression and executes the block or single statement that follows. The parenthesis are required.

If you only need a single statement, you can use the form that does not define a block. However, it never hurts to define the block and can save time later.

The if else Statement

- The *if else* statement examines an arbitrarily complex boolean expression

```
if ( <expr> )
{
    // statements
}
else
{
    // statements
}
```

5 – 5

The "if – else" statement lets you conditionally execute one of two blocks. And as shown on the previous page, if you have only one statement in the block, you can omit the curly braces.

The if else if Statement

- The *if else if* statement examines an arbitrarily complex boolean expression

```
if ( <expr> )
{
    // statements
}
else if ( <expr> )
{
    // statements
}
else
{
    // statements
}
```

5 – 6

The "if – else if – else" lets you create complex decisions. Note that the final "else" is optional.

The if else if Statement, cont'd

```
1  public class FlowControl
2  {
3      public static void main (String[] args)
4      {
5          java.util.Random rand = new java.util.Random();
6          int a = rand.nextInt();
7          int b = rand.nextInt();
8          int c = rand.nextInt();
9          if ( (a + b) > c )
10              System.out.println ( "Sum is greater" );
11          else if ( (a+b) == c )
12              System.out.println ( "The same" );
13          else
14              System.out.println ( "Sum NOT greater" );
15      }
16  }
```

5 – 7

This simple Java program retrieves three random integers, then determines if the sum of the first two is greater than, less than, or equal to the third.

Quick Practice

- The *System.in.read()* method accepts a keystroke from the user as a *char*. Write a program fragment that accepts a keystroke and calls the *System.exit(0)* method if the keystroke is an upper or lower case Q. For any other keystroke, display the keystroke character to the console.

```
char key = (char)System.in.read();
```

5 – 8

The while Loop

- The *while* loop evaluates a boolean expression and stays in the loop as long as the express is *true*

```
while ( <expr> )
{
    // statements
}

do
{
    // statements
}
while ( <expr> );
```

5 – 9

The difference between "while" and "do while" is that "do while" always does at least one pass through the loop.

Also note that if you have only one statement, you can omit the curly braces.

The while Loop, cont'd

```
1  public class FlowControl
2  {
3      public static void main (String[] args)
4      {
5          int i = 0;
6          while ( i < 10 )
7          {
8              System.out.println ( "in while: " + i );
9              i++;
10         }
11         do
12         {
13             System.out.println ( "in do-while: " + i );
14             i++;
15         } while ( i < 10 );
16     }
17 }
```

5 – 10

This Java program has two loops. The first, from lines 6 to 10, runs 10 times with "i" ranging from 0 to 9.

The second loop, from lines 11 to 15, runs once.

The for Loop

- The *for* loop is a convenient and concise form of the *while* loop

```
for ( <init statement(s)>;<expr>;<exec statement(s)> )  
{  
    // statements  
}
```

5 – 11

The "for" loop is "syntactic sugar" that lets you concisely write common looping constructs. The syntax requires a parenthesis with three fields, separated with semicolons. The first field contains statements that are executed before the loop begins -- if you have multiple statements, separate them with commas. The second field (between the semicolons), must be an expression that evaluates to a boolean -- the loop continues as long as the expression is true. The third field is any statements that are executed at the end of the loop -- again, if you have more than one statement, use commas to separate them.

The for Loop, cont'd

```
1  public class FlowControl
2  {
3      public static void main (String[] args)
4      {
5          for ( int i = 0; i < 10; i++ )
6          {
7              System.out.println ( "for loop: " + i );
8          }
9          for ( int i = 10,j = 0; i > 0; i--, j++ )
10         {
11             System.out.println ( "(" + i + "," +
12                                 j + ")" );
13         }
14     }
15 }
```

5 – 12

Here we show two examples of the "for" loop. The first, is a simple loop that counts from 0 to 9. The second illustrates using multiple "initial statements" and "executing statements". Its output looks like: (10,0)
 (9,1)
 (8,2)
 (7,3)
 (6,4)
 (5,5)
 (4,6)
 (3,7)
 (2,8)
 (1,9)

The for Loop, cont'd

```
1  public class FlowControl
2  {
3      public static void main (String[] args)
4      {
5          int i = 0;
6          for ( ; i < 10; )
7          {
8              System.out.println ( "for loop: " + i );
9              i++;
10         }
11     }
12 }
```

5 – 13

This "for" loop works the same as a "while" loop shown a few pages ago. It illustrates that you can write any "for" loop as a "while" loop and also shows that you can leave portions of the "for" loop empty, with semicolons as placeholders.

Breaking Out of a Loop

- Sometimes it's useful to exit a loop prematurely using the *break* statement

```
1  public class FlowControl
2  {
3      public static void main (String[] args)
4      {
5          java.util.Random rand = new java.util.Random();
6          while ( true )
7          {
8              int i = rand.nextInt();
9              if ( i < 0 )
10                 break;
11              System.out.println ( i );
12          }
13      }
14  }
```

5 – 14

The switch-case Statement

- The *switch-case* statement lets you choose between values of an integer or character type

```
switch ( <int or char type> )
{
    case <int or char const>:
        // statements
        break;
    case <int or char const>:
        // statements
        break;
    default:
        // statements
        break;
}
```

5 – 15

The "switch-case" statement has some rather odd characteristics, but is sometimes useful. The syntax requires parenthesis around an integer or character variable and then zero or more "cases" and an optional "default". At the end of each case, you can use the "break" statement to terminate the case -- if you omit it, then execution falls through to the next case.

The allowable types for the <int or char type> are "byte", "char", "short" or "int".

The switch-case Statement, cont'd

```
1  public class FlowControl
2  {
3      public static void main (String[] args)
4      {
5          java.util.Random rand = new java.util.Random();
6          int a = rand.nextInt();
7          switch ( a )
8          {
9              case 1:
10                 System.out.println ( "Value is One!" );
11                 break;
12
13             case 42:
14                 System.out.println ( "Value is 42!" );
15                 break;
16
17             default:
18                 System.out.println ( "Value is something else" );
19         }
20     }
21 }
```

5 – 16

This examples uses the "switch-case" statement to determine if a random integer has a value of 1, 42 or something else.

The switch-case Statement, cont'd

```
1  public class FlowControl
2  {
3      public static void main (String[] args)
4          throws Exception
5      {
6          while ( true )
7          {
8              char c = (char)System.in.read();
9              switch ( c )
10             {
11                 case 'd':
12                 case 'D':
13                     System.out.println(new java.util.Date());
14                     break;
15                 case 'q':
16                 case 'Q':
17                     System.exit ( 0 );
18             }
19         }
20     }
21 }
```

5 – 17

Here's another "switch-case" example that illustrates a reasonable use of the case "fall-through".

Quick Practice

- Write a loop that counts from 10 to 100 by fives and displays each value:

5 – 18

Chapter Summary

In this chapter, you learned:

- The basics of Java conditionals
- About looping constructs

Introduction to Object Orientation

- What is an Object?
- Three Pillars of OO

6 - 1

What is Object-Oriented Programming?

- Object oriented programming (OO) is a programming philosophy
- OO is also a set of techniques to achieve that philosophy
- Some programming languages, like Java, provide support for OO techniques
- OO will not solve all of the worlds programming ills

6 – 2

Object-Oriented Languages

- To best apply OO techniques, you should use a language written expressly for OO
- OO languages include:
 - Smalltalk
 - Java
 - C#
 - C++

6 – 3

Goals of Object Orientation

- Reuse (the Holy Grail)
- Easier to understand code
- Malleability
- Easier to maintain and debug
- Manage complexity for large projects
- Model the real world

6 – 4

The Three Pillars

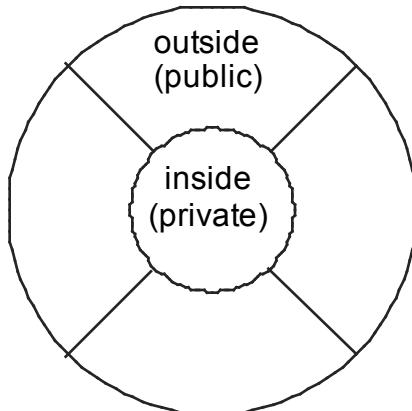
- To be considered a "true" OO language, the language should support:
 - Encapsulation
 - Inheritance
 - Polymorphism

6 – 5

Peter Wegner of Brown University devised a set of requirements for languages, known as the Wegner Definitions. In the Wegner Definitions, to be considered an object-oriented language, the language must support objects, inheritance and polymorphism.

What is an Object?

- An object is a self-contained chunk of state (data) and behavior (functions)
- Objects have private and public sections
- You can think of objects as being "alive", e.g. we might say that an object does this or that (executes its behavior)



6 - 6

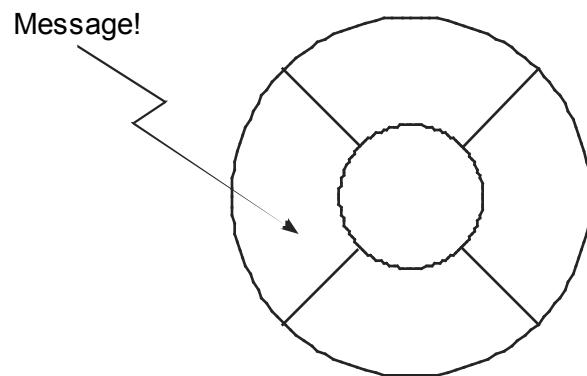
Why are Objects Useful?

- People naturally think of objects in the "real world"
- We intuitively understand how real-world objects have state (e.g. a dog's size) and behavior (e.g. eat)
- Often, we can model real-world objects with objects in a program (e.g. a BankAccount object)

6 – 7

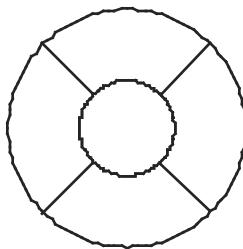
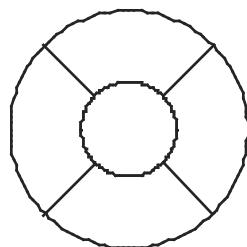
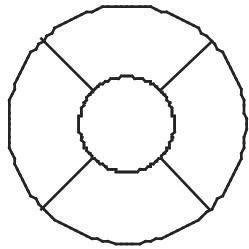
Objects Respond to Messages

- We refer to executing an object's behavior as "sending the object a message". In Java, we "send messages" by calling methods in the object
- The methods then typically operate on the state (data)



6 - 8

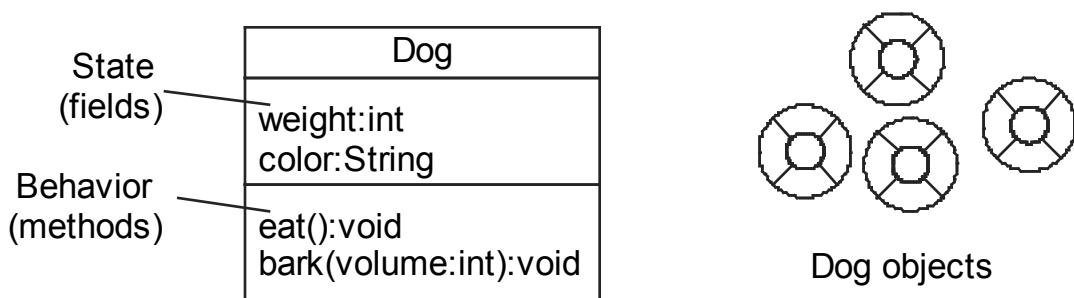
Objects and Classes



6 - 9

A Class is Like a Cookie Cutter

- In most O-O languages, you must define the class before you can create objects
- You can then use the class to "stamp out" objects of that class



Writing a Class

- A class is an abstraction that defines a collection of data (state) and the functions (behavior) that operate on the data

```
public class Dog
{
    // define fields (state)

    // define methods (behavior)
}
```

6 – 11

In Java, by convention, class names begin with an upper-case character.

Sample Java Class

6 – 12

You will write the class for Dog during lecture.

Instantiating Objects

- A class is like a cookie-cutter for objects
- The act of creating objects is referred to as *instantiation*

```
1  public class TestDog
2  {
3      public static void main(String[] args)
4      {
5          Dog lassie = new Dog();
6          Dog fido = new Dog();
7          lassie.eat();
8          fido.bark(8);
9      }
10 }
```

6 – 13

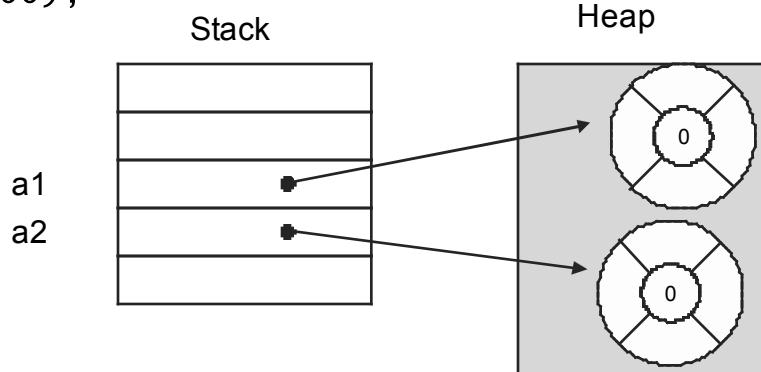
The Java "new" operator instantiates objects.

Since Java supports runtime garbage collection, there is no "delete" operator -- the Java Virtual Machine automatically gets rid of unreferenced objects.

Methods Operate on Objects

- Methods typically manipulate the object's state
- Methods invoked on one object don't affect other objects

```
1 Account a1 = new Account();  
2 Account a2 = new Account();  
3  
4 a1.withdraw (100);  
5 a2.deposit (300);
```



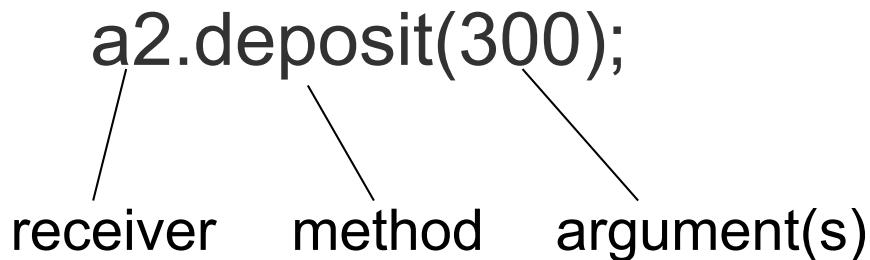
6 - 14

This is a very important notion in regards to object-orientation. You should think of methods as being like messages sent to an object -- messages processed by one object have no effect on other objects of the same type.

Methods Must Have a Receiver

- To invoke a method, you must explicitly or implicitly specify the **receiver** of the method
- This makes Java different than languages like C or C++ where you can have global methods
- The receiver is normally an object, but is a class for **static** methods (covered later)

```
Account a2 = new Account();
```



6 – 15

This is why referring to calling methods as "sending a message" is useful -- it emphasizes that the method must run in the context of an object (or a class, in the case of static methods).

Implicit Receivers

- If you call a method from within the same class, you can either explicitly specify **this** as the receiver, or use it implicitly

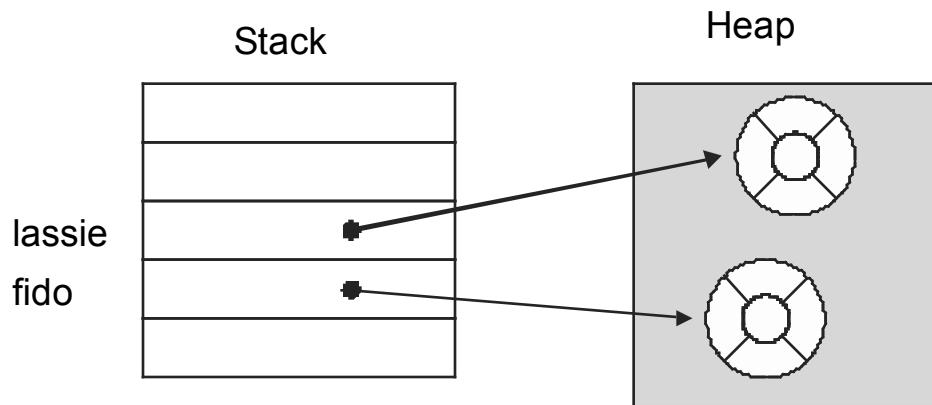
```
1  public class Test
2  {
3      public void method1()
4      {
5          }
6      public void method2()
7      {
8          this.method1();
9          method1();
10     }
11 }
```

6 – 16

The keyword "this" means "the current object". In SmallTalk, the same concept uses the keyword "self", which might be more intuitive than Java's "this".

The Stack and the Heap

- The *new* operator carves out storage for the object from a memory area referred to as the heap



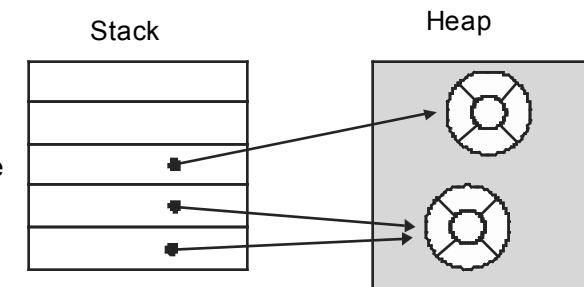
6 – 17

The heap is the free storage remaining in the Java Virtual Machine.

Comparing References

- You can use the *equality* operator to compare references

```
1 Dog lassie = new Dog();  
2 Dog fido = new Dog();  
3 Dog spike = fido;  
4  
5 if (fido == spike)  
6     System.out.println ("Same!");  
7 if (lassie == fido)  
8     System.out.println ("Same!");
```



6 - 18

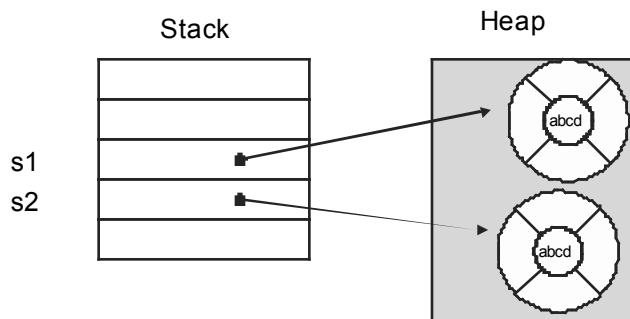
When you use the equality operator on references, it compares the reference values, not the objects themselves.

Comparing References, cont'd

- When applied to references, the *equality* operator compares the references themselves, not the objects to which the references refer

```
1  String s1 = new String("abcd");
2  String s2 = new String("abcd");
3
4  if (s1 == s2)
5      System.out.println("Same!");
6  if (s1.equals(s2))
7      System.out.println("Equals!");
```

6 – 19



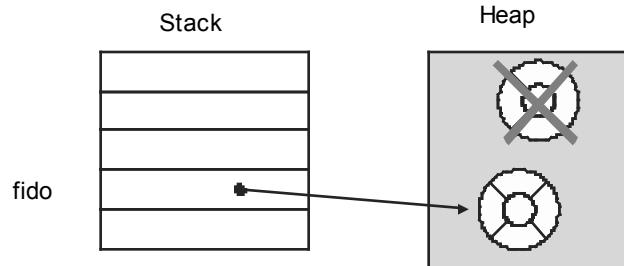
Even though the two String objects have the exact same characters, the equality operator applied to the objects' references returns false. That's because there are two distinct objects on the heap, each with a different address.

To compare the contents of the objects, the String class provides the equals() method, which returns true only if the two Strings have the same characters (it is a case-sensitive comparison).

Garbage Collection

- The garbage collector runs as a *background thread* that periodically scans the stack and the heap looking for unreferenced objects
- The garbage collector automatically deletes all unreferenced objects

```
{  
    Dog lassie = new Dog();  
    Dog fido = new Dog();  
    . . .  
    lassie = null;  
    . . .  
}
```



6 – 20

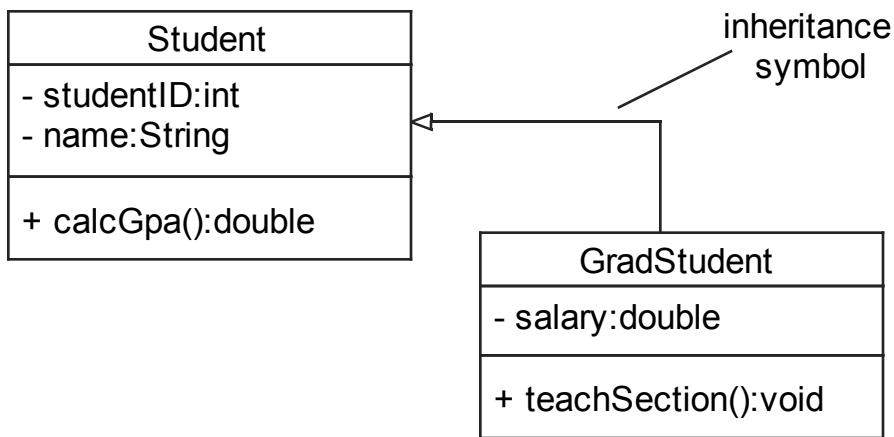
Java defines the keyword "null" which you can assign to a reference to indicate that the reference is no longer in use.

An object on the heap becomes unreferenced either when you explicitly set all of its references to null, or if the local variables that contain the object's references go out of scope.

You can force the garbage collector to run by calling the `gc()` method on the `System` class, which is part of the Java class library included in the JRE.

Introduction to Inheritance

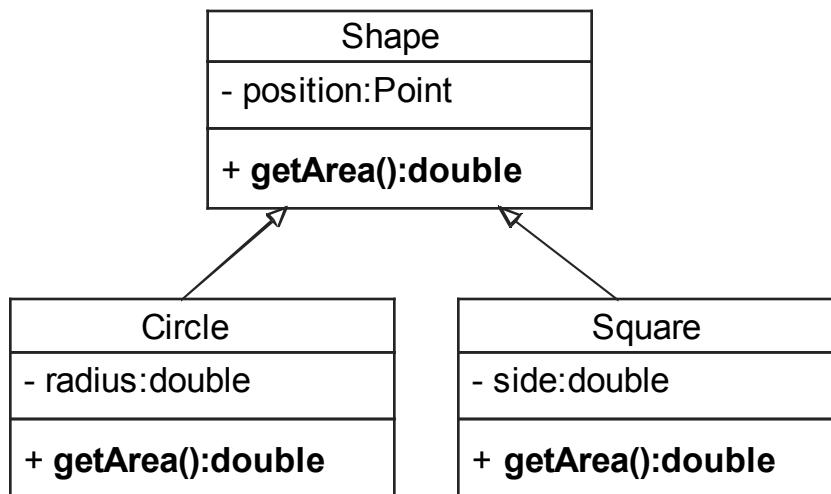
- A subclass inherits state and behavior from its superclass
- Inheritance makes it easier to extend and reuse code
- **Important:** A subclass object is a superset of all of its superclasses



6 – 21

Introduction to Polymorphism

- Polymorphism means "many forms"
- Polymorphism helps you design extensible code where subclasses override behavior from superclasses



6 – 22

Synonym Alert!

Generic OO Term	C++ Term	Java Term
state	member data	field
behavior	member function	method
inherit from (specialize)	subclass	extend
superclass	base class	superclass
subclass	derived class	subclass

6 – 23

Introduction to Packages

- Java lets you organize your classes into *packages*
- If a class needs to use a class from another package, it must either fully-qualify the reference or use the *import* statement

```
1  package mypackage;
2  import java.util.Date;
3  import java.net.*;
4
5  public class Test
6  {
7      public static void main(String[] args)
8      {
9          Date d = new Date();
10         Socket s = new Socket();
11         java.io.File myFile =
12             new java.io.File ("test.txt");
13     }
14 }
```

6 – 24

Packages define namespaces so you can potentially re-use class names (i.e. you can use the same class name as long as the classes are in different packages). You can also nest packages within other packages.

If present, the "package" statement must be the first non-comment line in the file.

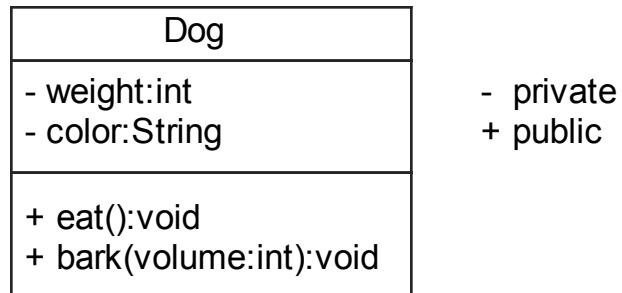
We refer to classes that have no package statement as being in the "unnamed" or "default" package.

The compiler automatically imports everything from the `java.lang` package for you.

Note that the import statement has two syntaxes: one that imports a single type and another that imports all of the types from a specified package.

Encapsulation

- OO languages enforce encapsulation by restricting access to class functions and data
- Encapsulation makes it easier to maintain and reuse code (black box approach): you can safely modify hidden parts



6 – 25

Using Encapsulation

```
1  public class Dog
2  {
3      private int weight;
4      private String color;
5
6      public void eat()
7      {
8          }
9
10     public void bark (int volume)
11     {
12         }
13 }
```

6 – 26

Java Encapsulation Summary

```
public class Test
{
    public int a;
    private int b;
    protected int c;
    int d;
}
```

Access Specifier	Within same class	Within a subclass	Within same package	Outside world
private	●			
protected	●	●	●	
default	●		●	
public	●	●	●	●

6 - 27

Java provides four ways to encapsulate fields and methods: public, private, protected and default, which is known as "package-level".

Get/Set Methods

- Since fields are normally private, you can provide controlled access with get/set methods

```
public class Account
{
    private double balance = 100.00;
    .
    .
    public double getBalance()
    {
        return this.balance;
    }
    public void setBalance(double balance)
    {
        this.balance = balance;
    }
    .
    .
}
```

6 – 28

Get/Set methods, also known as accessor/mutator methods, give the outside world controlled access to private fields. In other words, the only way code outside of the class can read or write private fields is via the get/set methods, in which we can write code to do whatever we want (e.g. write a "set" method that only allows changes within a range of values). If you were to make fields public, then you would have no control over how the fields are read and written.

This code also shows using the Java "this" keyword, which refers to the current object. If you omit this keyword on a field access, the compiler assumes it. Note that for the "set" method, we need to use it to differentiate between the name of the field and the name of the parameter to the "set" method.

Get/Set Methods, cont'd

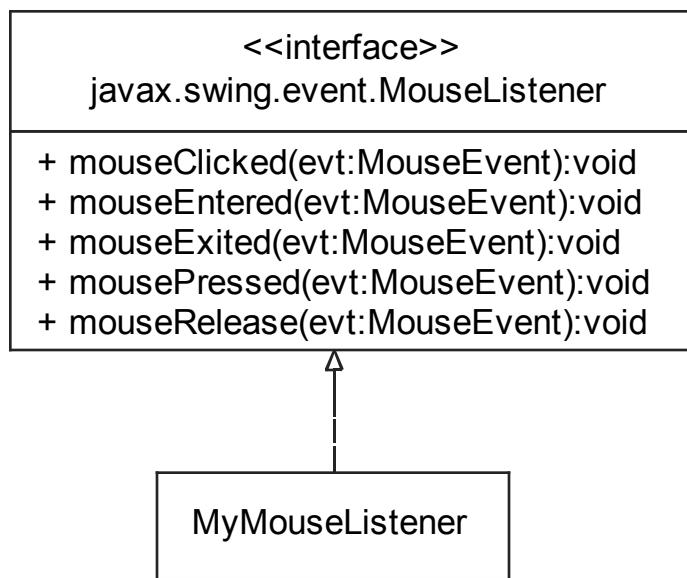
```
1  public class TestAccount
2  {
3      public static void main(String[] args)
4      {
5          Account a1 = new Account();
6          Account a2 = new Account();
7
8          a1.balance = 100; // ERROR!
9
10         a1.setBalance(100);
11         a2.setBalance(200);
12
13         double d1 = a1.getBalance();
14         double d2 = a2.getBalance();
15     }
16 }
```

6 – 29

This code demonstrates creating a couple of Account objects and then exercising the get/set methods.

Introduction to Interfaces

- An **interface** is a group of methods with no implementation
- Classes **implement** interfaces by providing the methods defined in the interface



6 – 30

The Java Runtime Environment defines many interfaces that define a set of behavior that a class can provide. Here we show the `MouseListener` interface, which specifies the methods that a Java GUI component calls when the mouse interacts with the component.

You can also create your own interfaces to define sets of behavior.

Quick Practice

- Write a Java class named `Policy` that represents an insurance policy. The class should maintain state for the policy holder's name, the policy amount and whether the holder is a smoker. The class should define a behavior named `rate` that returns true if the holder is a non-smoker and the amount is less than 100000 and false otherwise.
- Write another class named `TestPolicy` with a *main* method that creates an instance of the above class and prints the object's "rating" to the console.

6 – 31

Chapter Summary

In this chapter, you learned:

- The basics of object orientation
- About the three pillars:
encapsulation, inheritance and
polymorphism

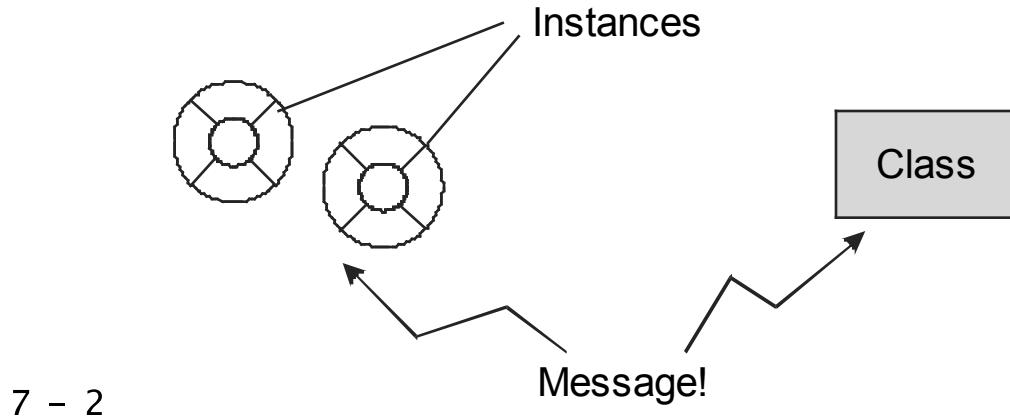
Methods

- Instance Methods
- Constructors
- Static Methods

7 - 1

What is a Method?

- Methods represent behaviors in an object or a class
- **Instance** methods run as part of an individual object, while **static** methods run as part of the class
- Calling a method is sometimes referred to as "sending a message"



7 - 2

Recall that an object consists of state and behavior -- in Java, we write methods to provide behaviors. In this chapter, we will start by looking at instance methods and later cover static methods.

Calling a Method

- To call an **instance** method, you must first create a object, then use the object reference to invoke the method
- To call a **static** method, you invoke the method on the class itself

General Form

```
type returnVal = receiver.method(arg1, arg2)
```

```
1 public class Test
2 {
3     public static void main(String[] args)
4     {
5         java.util.Random rand = new Random();
6         int i = rand.nextInt();    // instance method
7
8         double result =
9             Math.sqrt(8.33);    // static method
10    }
11 }
```

7 - 3

Method Receivers

- Method invocations must **always** be directed to a **receiver**, which is either an object (for instance methods) or a class (for static methods)
- In some cases, the receiver can be implicit

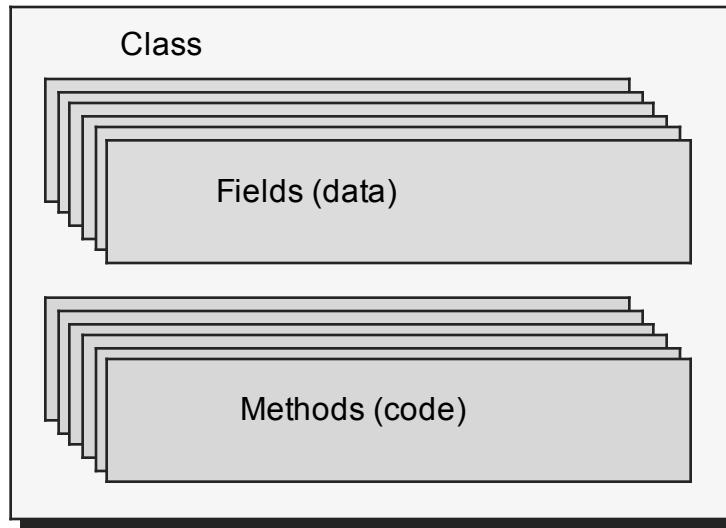
```
1  public class Test
2  {
3      public void myMethod()
4      {
5          anotherMethod();
6          this.anotherMethod();
7      }
8      public void anotherMethod()
9      {
10     }
11 }
```

7 – 4

If a method in a class needs to call another instance method in the same class, you can use the "this" keyword to explicitly specify the receiver.

Class Structure

- In Java, all code must be contained within a method (a.k.a function or subroutine)
- In turn, methods must reside within a class



7 – 5

The "class" is the basic building block of Java -- each class typically defines methods (behaviors) and fields (state). In this chapter, we'll focus on the methods.

Method Syntax

- To define a method, you must specify the method's return type, name and arguments (if any)
- If the method has no return value, use *void*
- If the method has no arguments, you still must use parenthesis

```
[options] <return-type>
methodName ( [arg-type argName]* )
```

7 – 6

By convention, Java method names begin with a lower-case first character.

Instance Method Examples

```
1  public class Methods
2  {
3      public void method1()
4      {
5          // statements
6      }
7      public int method2 ( int x )
8      {
9          return x * x * x;
10     }
11     public double method3(int x, int y)
12     {
13         return Math.PI * Math.PI + x + y;
14     }
15 }
```

7 – 7

Here we show a class with three methods. The first method has no arguments and no return value. The second accepts a single integer and returns an integer. The third accepts two integers and returns a double.

Calling Instance Methods

- Unless a method is defined as *static*, you must create an object before calling the method
- These non-static, or *instance* methods run in the context of the object in which they reside

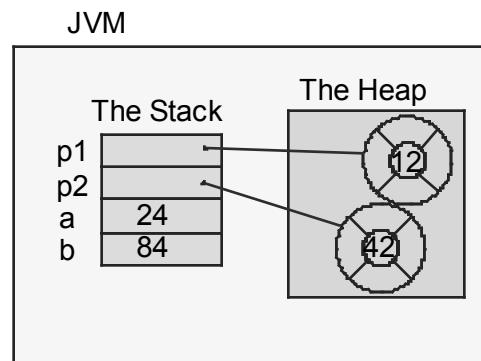
```
1  public class Person
2  {
3      private int age = 14;
4
5      public void setAge(int a)
6      {
7          age = a;
8      }
9      public int method1 (int z)
10     {
11         return age * z;
12     }
13 }
```

7 – 8

If you omit the "static" option when you define a method, the method is considered to be an "instance" method, which means it runs as part of an individual object. Instance methods typically operate on the object's state.

Calling Methods, cont'd

```
1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          Person p1 = new Person();
6          Person p2 = new Person();
7          p1.setAge(12);
8          p2.setAge(42);
9          int a = p1.method1(2);
10         int b = p2.method1(2);
11     }
```



7 – 9

Here we show a program that uses the class defined on the last page. It creates two objects and then calls instance methods on each. Note how the methods operate on the each individual object's data.

Quick Practice

- Given the *Methods* class shown on page 7, write a program fragment that creates an instance of the *Methods* class and calls each of its methods

7 – 10

Quick Practice 2

- Look up the *java.util.Random* class in the JavaDocs and write a program fragment that creates an instance of the *Random* class and calls its *setSeed* and *nextLong* methods.

Overloading Methods

- You can define multiple methods with the same name as long as each has a different *signature*
- This technique is referred to as *overloading*

```
1  public class Account
2  {
3      public void withdraw ( int amount )
4      {
5          // statements
6      }
7      public void withdraw ( double amount )
8      {
9          // statements
10     }
11 }
```

7 – 12

Overloaded methods with a class must have different signatures. Note that when determining a method signature, Java doesn't use the return value -- in other words, the number and/or types of the parameters must be different.

Calling Overloading Methods

- When a program calls an overloaded method, the compiler looks for a matching signature

```
1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          Account a = new Account();
6
7          a.withdraw ( 20 );
8          a.withdraw ( 55.55 );
9      }
10 }
```

7 – 13

If the compiler doesn't find a method with the correct name and signature, you get a compile error!

Constructor Methods

- Constructor methods are special methods that participate in the construction of objects
- Constructor methods must have the same name as the class and have no return value

```
1  public class Account
2  {
3      private double balance;
4
5      public Account()
6      {
7          // statements
8      }
9 }
```

7 – 14

The job of a constructor is to initialize the object so it's ready for use. The Java virtual machine calls constructors during the execution of the "new" operator.

Normally, constructors are public.

The new Operator

- The *new* operator does the following:
 - Allocates storage from the heap
 - Sweeps the storage to value zero
 - Runs field initializers
 - Calls a constructor
 - Returns the reference

7 – 15

Constructors run as part of the "new" operator.

Overloaded Constructors

- A class can define multiple constructors, as long as each has a different signature

```
1  public class Account
2  {
3      private int accountID;
4      private double balance;
5
6      public Account ( int id )
7      {
8          accountID = id;
9          balance = 0;
10     }
11     public Account ( int id, double bal )
12     {
13         accountID = id;
14         balance = bal;
15     }
16 }
```

Overloaded constructor methods have the same signature uniqueness rules as non-constructor methods.

Often, we use overloaded constructors to provide different ways to initialize new objects.

Calling Constructors

- Calling programs never directly invoke a constructor
- Instead, the Java virtual machine calls a constructor during the *new* operator

```
1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          Account a1 = new Account ( 12 );
6          Account a2 = new Account ( 14, 45.67 );
7      }
8  }
```

7 – 17

When a calling program uses the "new" operator, Java looks for a constructor with a matching signature. If one does not exist, you get a compile error.

Note that by providing multiple constructors, the class author gives the class user flexibility in how they create objects.

Quick Practice

- Write a class named **Claim** that represents a claim on a life insurance policy. The class should have state for the date of the claim, the claim number and claim amount. Write two constructors in the class: one that accepts claim number and amount, and another that accepts data corresponding to all of the fields.
- Write another class named **TestClaim** with a *main* method that creates two objects, using both constructors.

The No-Argument Constructor

- If a class defines no constructors, Java provides one that accepts no parameters
- But if the class defines any constructors, Java no longer provides the no-argument constructor

Why will this code not compile?

```
1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          Account a1 = new Account ( 12 );
6          Account a2 = new Account ( 14, 45.67 );
7          Account a3 = new Account();
8      }
9  }
```

7 – 19

Assuming we are using the Account class defined a couple of pages ago, this code will get a compile error in line 7. That's because the Account class doesn't define a no-argument constructor.

The no-argument constructor is sometimes referred to as the "default" constructor since Java provides it if the class has no other constructors.

Calling One Constructor From Another

- To avoid duplicated code, a constructor can use the *this* operator to call another constructor

```
1  public class Account
2  {
3      private int accountID;
4      private double balance;
5
6      public Account()
7      {
8          // complex calculation here
9      }
10     public Account ( int id )
11     {
12         this();
13         accountID = id;
14     }
15 }
```

7 - 20

The "this" keyword refers to the current object. Here we are using it to invoke a constructor method.

Note that any call to another constructor must be the first executable line in the calling constructor.

Field Initializers

- If you wish to initialize fields to values known at compile time, you don't need to write a constructor

```
1  public class Account
2  {
3      private int accountID = 14;
4      private double balance = 100.00;
5      private String uniqueHashCode;
6
7      public Account()
8      {
9          // calculate unique hash code
10     }
11 }
```

7 – 21

If you don't initialize fields, the "new" operator sweeps them to "natural zero", which means zero for numeric types like integers and sets object references to null.

Note that you may still need to write a constructors to initialize fields to values not known at compile time (e.g. data from a database) or if the initialization takes more than a single assignment (e.g. you need to calculate a value).

Method Modifiers

- When you define a method, you can specify optional modifiers that configure the method

public
private
protected

Access to the method

final

Method cannot be overridden by a subclass

abstract

Method has no body (must be provided by a subclass)

static

Method is part of the class, not part of any individual object

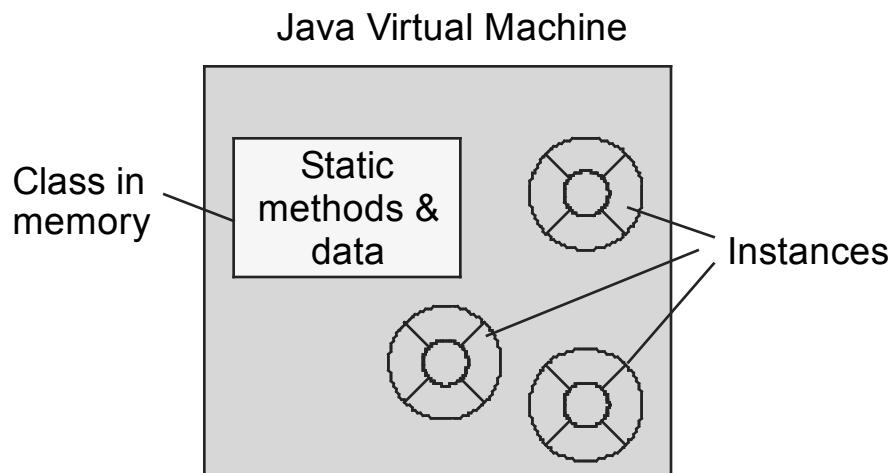
7 – 22

The "public" access modifier means that any code can call the method. The "private" modifier means that only other methods in the same class can call it. The "protected" modifier means that the method is also available to be called from subclasses.

If you don't specify any access modifier, the method has "package" protection, which means it can be called from other classes that are in the same package.

Static Methods

- **Static methods** are part of the class in which they are defined, not part of any individual object
- You can call static methods without creating a class instance
- Static methods cannot access non-static fields



7 – 23

Remember that instances are created by the "new" operator.

Static Methods, cont'd

```
1  public class MyClass
2  {
3      private static int x;
4      private int y;
5      public static void myMethod()
6      {
7          System.out.println (x); //OK
8          System.out.println (y); //ERROR
9      }
10 }
```

```
1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          MyClass.myMethod();
6      }
7 }
```

7 – 24

Implicit Receiver for Static Method

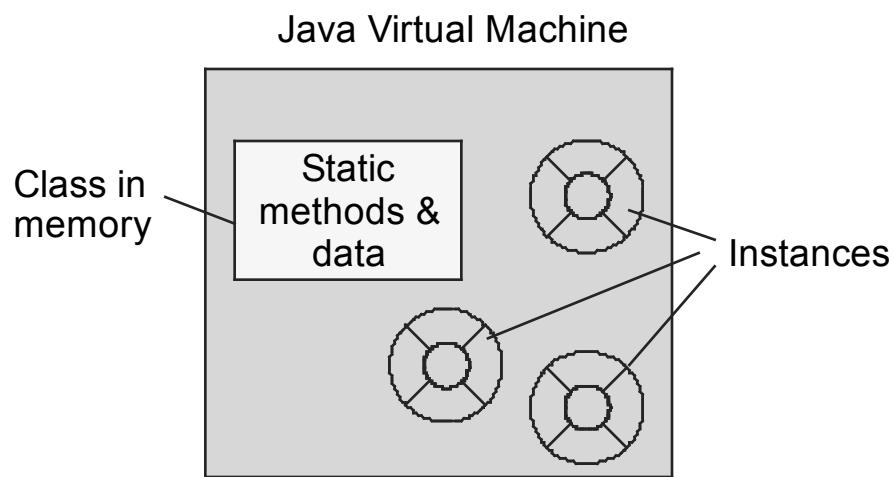
- If a method needs to call a static method in the same class, you can omit the class name as the receiver

```
1  public class TestQuadratic
2  {
3      public static double solveQuadraticPlus(double a,
4          double b, double c)
5      {
6          return (-b + Math.sqrt(b*b-4*a*c))/2*a;
7      }
8      public static void main(String[] args)
9      {
10         double result = solveQuadraticPlus(1,3,-4);
11         result =
12             TestQuadratic.solveQuadraticPlus(1,3,-4);
13     }
14 }
```

7 – 25

Static Fields

- Static **fields** are part of the class in which they are defined, not part of any individual object
- You can access static fields without creating a class instance
- Static fields can be accessed by either static or non-static methods



Static Fields, cont'd

```
1  public class MyClass
2  {
3      public static int x;
4  }
```

```
1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          System.out.println (MyClass.x);
6      }
7  }
```

Quick Practice

- Use the JavaDoc to look up the *java.io.File* class and its *createTempFile* method that accepts two Strings. Write a program fragment that creates a temporary file with prefix="myfile" and a *null* suffix. Then display the temporary file's *path* to the console.

```
1 import java.io.File;  
2  
3  
4 File f =
```

The final Keyword

- Final **methods** cannot be overridden in a subclass
- Final **fields** are read-only after construction

```
1  public class MyClass
2  {
3      public static final int count=12;
4
5      public final void myMethod()
6      {
7      }
8  }

1  public class Test
2  {
3      public static void main(String[] args)
4      {
5          System.out.println (MyClass.count);
6      }
7  }
```

Using static and final together lets you define constants. Note that when you define a final field, you must initialize it as part of the definition or in a constructor.

By convention, constants are written in all uppercase characters.

Also be aware that J2SE 1.5 defines enumerations, which are a more powerful way to create constants.

NOTE: You can also define classes as final -- such classes cannot be subclassed.

Chapter Summary

In this chapter, you learned:

- How to define methods
- About constructors and static methods

Exception Handling

- Traditional Error Handling
- Java Exceptions

8 - 1

Traditional Error Handling

- Traditionally, programmers have handled potential errors by checking return codes from functions

```
1  public void myMethod()
2  {
3      int errcode=someAPI();
4      if (errcode != 0)
5          processError(errcode);
6      else
7      {
8          // do other stuff
9      }
10 }
```

8 – 2

Traditional Error Handling Issues

- Error checking clutters the code
- In a fully debugged program, error checking adds unneeded overhead
- It's difficult to handle errors globally, especially given nested function calls

8 – 3

Is there a better way?

Java Exception Handling

- Instead of explicitly checking for errors, use a *try-catch* block

```
1  public void myMethod()
2  {
3      try
4      {
5          someAPI();
6          // do other stuff
7      }
8      catch ( SomeException exc )
9      {
10         processError(exc);
11     }
12 }
```

8 – 4

If no exception occurs within the "try" block, it executes normally and then skips over the "catch" block. If the "SomeException" occurs, then Java jumps out of the "try" block into the "catch" block -- this is where you can respond to the exceptional condition.

Advantages of Exception Handling

- Code that incurs no exceptions doesn't need overhead of *if* statements checking for errors
- You can write multiple *try–catch* blocks within a single method
- Compiler will force you to handle exceptions if you call methods that are marked as throwing an exception
- One downside is that the syntax is rather verbose

Uncaught Exceptions

- If the Java virtual machine detects an uncaught exception, it terminates the program and displays a stack trace

```
1  public class Uncaught
2  {
3      public static void main(String[] args)
4      {
5          myMethod();
6      }
7      public static void myMethod()
8      {
9          String s = null;
10         int len = s.length();
11     }
12 }
```

```
C>java Uncaught
java.lang.NullPointerException
at Uncaught.myMethod(Uncaught.java:10)
at Uncaught.main(Uncaught.java:5)
Exception in thread "main"
```

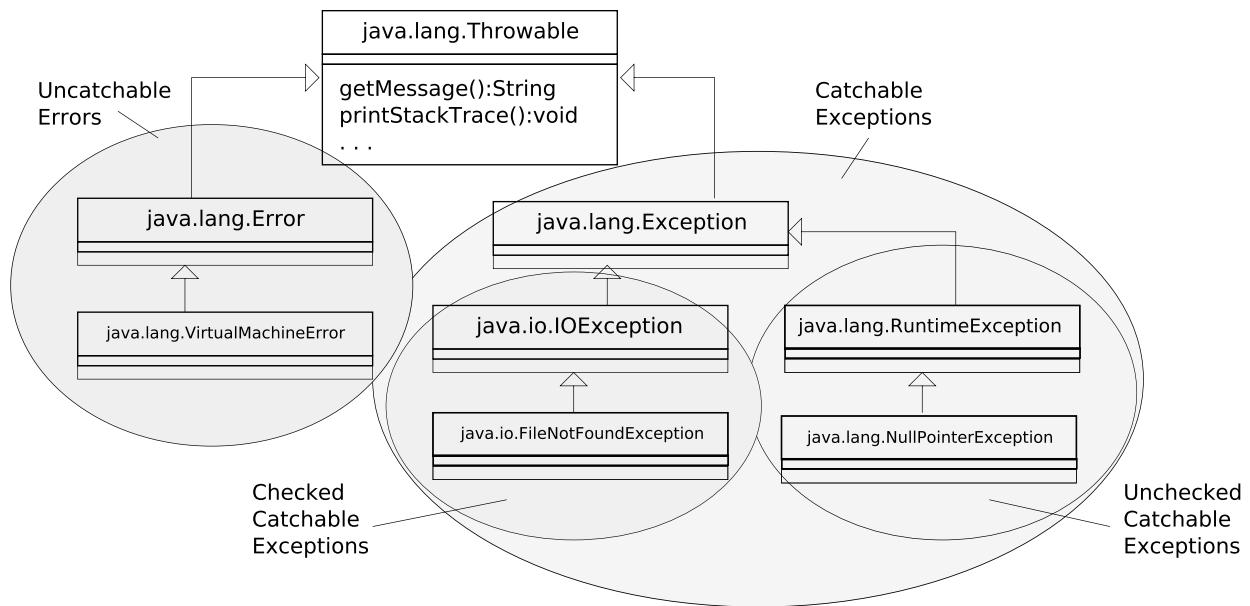
8 – 6

Uncaught exceptions are passed back through any nested method calls -- if nobody catches it, the exception finally makes it to the virtual machine.

Here we show a program that attempts to use a null reference, which causes the NullPointerException.

Exceptions Are Objects

- Java supports three categories of exceptional conditions: **errors**, **checked exceptions** and **unchecked exceptions**



8 – 7

Like all objects, exceptions are members of a class. In addition, exception classes are related by inheritance. `java.lang.Throwable` is the root of all exception classes.

Here we show some of the exception classes in the Java class library -- there are many, many more and you can create your own.

Errors are conditions that you cannot handle in your program, while Exceptions can be caught. Catchable exceptions themselves fall into two categories -- checked exceptions, which the compiler requires you to handle, and unchecked exceptions which you can handle, but are not required to.

When you are reading the JavaDoc for an exception, you can see whether it's checked or unchecked by looking for `java.lang.RuntimeException` in its inheritance hierarchy. If it's there, then the exception is unchecked.

Multiple Catch Blocks

```
1  try
2  {
3      FileReader fr =
4          new FileReader ( "/temp/test.txt" );
5      int ch = fr.read();
6      while ( ch != -1 )
7      {
8          System.out.print ( (char)ch );
9          ch = fr.read();
10     }
11     fr.close();
12 }
13 catch (FileNotFoundException e)
14 {
15     System.out.println ( "Unable to open file" );
16 }
17 catch (IOException e)
18 {
19     System.out.println ( "Error reading file" );
20 }
```

8 – 8

This code demonstrates using code in a "try" block that throws more than type of exception. The statement in line 3 throws the FileNotFoundException, while the code in lines 4, 8 and 10 can throw the generic IOException.

Handle the Most Specific Exception First

- When an exception occurs, the JVM uses the first matching *catch* block
- Since exception classes are related by inheritance, you must code any catch block before another catch block of its superclass

```
1  try {  
2  . . .  
3  }  
4  catch (FileNotFoundException e) {  
5  }  
6  catch (IOException e) {  
7  }  
8  catch (Exception e) {  
9  }
```

8 – 9

The code shown here demonstrates how to catch specific exceptions before the more generic ones. In this case, the "Exception" catch block acts like a catch-all for exceptions that we will not handle specifically.

Passing Exceptions Back

- Instead of writing a *try-catch* block, a method can pass exceptions back to its own caller
- This lets you decide whether to handle exceptions locally or at a more global level

```
1 public void readFile()
2     throws IOException,FileNotFoundException
3 {
4     FileReader fr =
5         new FileReader ( "/temp/test.txt" );
6     int ch = fr.read();
7     while ( ch != -1 )
8     {
9         System.out.print ( (char)ch );
10        ch = fr.read();
11    }
12    fr.close();
13 }
```

8 – 10

A method that calls other methods that are marked as throwing exceptions has two choices on how to handle the exceptions:

- 1.The calling method can use a *try-catch* block as shown previously
- 2.The calling method can pass the exceptions back to its own caller

The finally Block

```
1  public static void myMethod()
2  {
3      try
4      {
5          // statements
6          return;
7      }
8      catch ( Exception exc )
9      {
10         System.out.println ( exc.getMessage() );
11         return;
12     }
13     finally
14     {
15         // always executes!
16     }
17 }
```

8 – 11

The "finally" block is an optional part of the "try" syntax. If present, it must be after any "catch" blocks.

Even if the code in a "try" or "catch" block returns or throws an exception, the virtual machine guarantees it will execute the statements in the finally block.

This is especially useful in situations where you want to be sure to close a resource (e.g. a database connection).

Throwing Exceptions

- A method can raise, or **throw** one or more exceptions to indicate an exceptional condition
- Since exceptions are objects, we use the **new** keyword, specifying one of the exception class's constructors

```
1 // java.lang.IllegalArgumentException is unchecked
2 public static void myMethod2 (int x)
3 {
4     if ( x < 0 )
5         throw new IllegalArgumentException (
6             "I don't like negatives!" );
7
8     // statements
9 }
```

8 – 12

IllegalArgumentException is an unchecked exception since it's a subclass of java.lang.RuntimeException.

We create a new exception object so that we can throw it and that we pass a message string to the exception class's constructor. That "human-readable" string is stored by the java.lang.Throwable superclass.

Note that when a method throws an exception, that terminates the method's execution and returns control back to the caller.

If a method throws any CHECKED exceptions (i.e. not a subclass of RuntimeException, then it must list those in the "throws" clause of the method definition:

```
public static void myMethod3()
    throws java.io.FileNotFoundException
{
    . . .
    throw new FileNotFoundException("Uh oh");
}
```

Writing an Exception Class

- Java applications can define their own exception classes to represent domain-specific exceptions
- Your choice of superclass determines whether the new exception class is checked or unchecked

```
1  public class InsufficientFundsException
2      extends Exception
3  {
4      public InsufficientFundsException (
5          String msg )
6      {
7          super ( msg );
8      }
9  }
```

8 – 13

Here we've defined an exception for a banking application that an application can throw if a withdrawal would result in a negative balance. Since we choose `java.lang.Exception` as the superclass, this is a "checked" exception.

The "super" keyword calls a constructor in the `InsufficientFundsException`'s superclass (i.e. `Exception`). Note that writing a constructor in the exception class is optional.

Writing an Exception Class, cont'd

```
1  public class CheckingAccount
2  {
3      private double curBalance;
4
5      public void writeCheck ( double amt )
6          throws InsufficientFundsException
7      {
8          if ( curBalance - amt < 0 )
9              throw new InsufficientFundsException (
10                  "Out of money" );
11
12         curBalance -= amt;
13     }
14 }
```

8 – 14

Here we show using the exception class defined on the last page. Since the `InsufficientFundsException` is a "checked" exception, we list it in the method's definition.

Quick Practice

- Use the JavaDocs to look up the `java.net.Socket` class. Examine the constructor that accepts a hostname String and a "port" integer. Then write a program fragment that creates a Socket to connect to `google.com` on port 80. Use try/catch blocks as necessary.

Chapter Summary

In this chapter, you learned:

- Exception handling fundamentals
- How to handle exceptions
- Writing your own exception classes

Arrays and Collections

- Using Arrays
- Using Collections

9 - 1

What is an Array?

- Arrays are simple data structures that store homogeneous data
- You access arrays using an index, which provides fast random access to array elements
- Array indexes start at zero

```
1  int[] intArray = new int[10];
2  double[] doubArray = new double[5];
3
4  intArray[0] = 3;
5  intArray[1] = 77;
6  . . .
7  int x = intArray[1];
8
9  doubArray[0] = 77.77;
10  doubArray[1] = 6444.33;
11  . . .
```

9 – 2

Like objects, arrays must be allocated with the "new" operator before you can use them.

The "new" operator sets all of the array elements to "logical zero" for the type of the array. For example, integers are set to the value zero and arrays of object references are set to "null".

Using Arrays

- It's common to use a **for** loop to access array elements
- Remember to start the loop at zero!

```
1  public static void main (String[] args)
2  {
3      int[] intArray = new int[10];
4      for ( int i = 0; i < intArray.length; i++ )
5          intArray[i] = i * 2;
6  }
```

0	1	2	3	4	5	6	7	8	9
0	2	4	6	8	10	12	14	16	18

9 – 3

Here we demonstrate using a "for" loop to initialize array elements. Note that the loop starts at zero and the ending index is "n – 1", where "n" is the number of elements allocated for the array.

Array Bounds Checking

- The virtual machine checks each array access to ensure that the index is within bounds
- Violations generate the `ArrayIndexOutOfBoundsException`

```
for ( int i = 1; i <= intArray.length; i++ )  
    System.out.print (intArray[i]);
```

```
C>java Arrays  
24681012141618java.lang.ArrayIndexOutOfBoundsException: 10  
at introjava.UsingArrays.main(UsingArrays.java:12)  
Exception in thread "main"
```

9 – 4

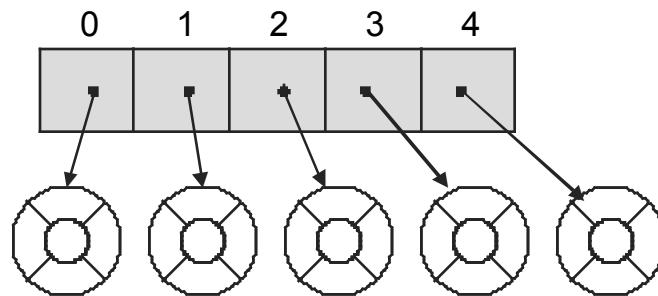
Here we show a common mistake with arrays -- indexes should start at zero and range to n-1 where n is the array's length.

Arrays of Object References

- When you allocate an Array that holds a reference type, the *new* operator sets the array's contents to *null*
- Before use, you must first set the references so they refer to actual objects

```
1  Student[] stArray = new Student[5];
2
3  for ( int i = 0; i < stArray.length; i++ )
4      stArray[i] = new Student();
```

9 – 5



Arrays of object references are quite useful, but you have to be careful to remember that allocating the array itself does not allocate the objects!

Initializing Arrays

- For types that have literal initializers, you can allocate and initialize an array in one step

```
1 int[] a1 = {45, 55, 77};  
2  
3 String[] a2 = {"ab", "def", "x" };
```

9 – 6

Defining and initializing an array in one step is a useful shortcut.

Quick Practice

- Write a program fragment that defines an array to hold 5 doubles. Then write a loop to initialize the doubles to multiples of Math.PI.

9 - 7

The Enhanced for Loop

- In Java 5 or later, you can use the **enhanced for** loop to iterate an array

```
1  public static void main (String[] args)
2  {
3      int[] intArray = new int[10];
4      for (int i = 0; i < intArray.length; i++)
5          intArray[i] = i * 2;
6
7      for(int i : intArray)
8      {
9          System.out.println(i);
10     }
11 }
```

9 – 8

You can also use the "enhanced for loop" on so-called generic collections.

Introduction to Collections

- Arrays are useful, but are difficult to work with if you don't know in advance how many elements the array will hold
- It's also difficult to re-size an array or to insert a new element in the middle of an array
- Java provides *collection classes* that address these issues
- **Note:** Collections store only object references, not primitive values

9 – 9

Most of the collection stuff is in the "java.util" package.

Collections Can Automatically Grow

- As you add elements to a collection, it will automatically resize itself
- Of course, re-sizing takes time!

```
1  java.util.ArrayList<String> myList =  
2      new java.util.ArrayList<String> ( 10 );  
3  
4  for ( int i = 0; i < 1000; i++ )  
5      myList.add ( "adding element " + i );
```

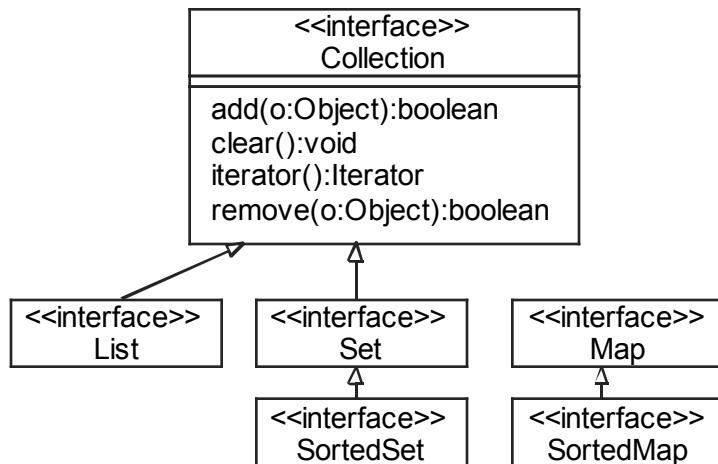
9 – 10

Here we show adding String objects to a collection.

You could conceivably write code to do this with arrays, but collections resize automatically!

Basic Collection Types

- The *Collection* interface is the base type for the collections in Java
- **Lists** are sequential collections that allow duplicate elements
- **Sets** reject duplicates and are optionally sorted
- **Maps** are key/value collections where the key must be unique



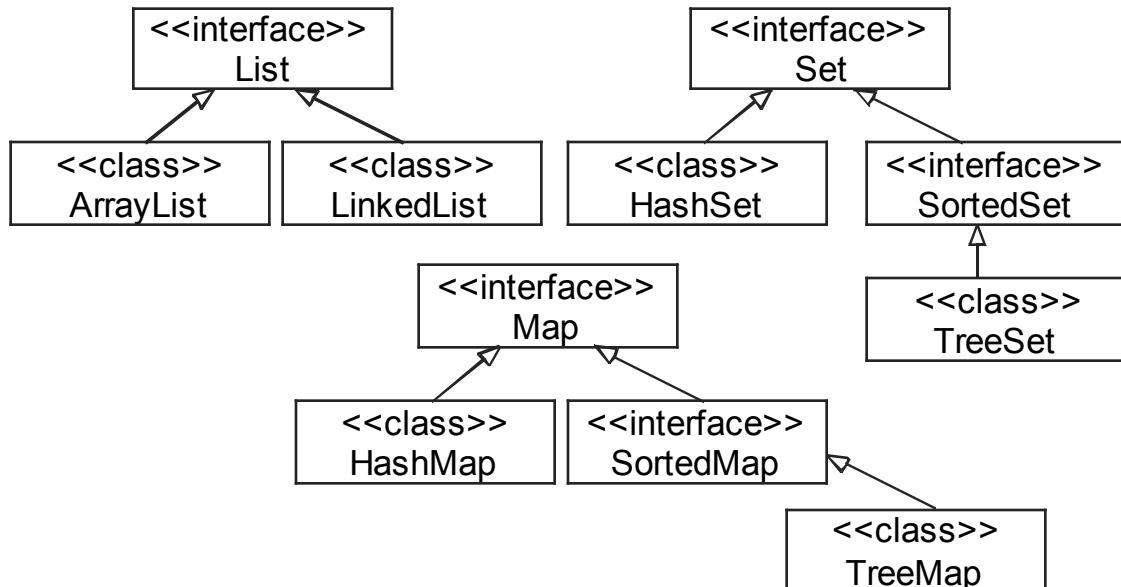
9 – 11

These are all "interfaces", which are like classes but cannot be instantiated. Interfaces define a set of methods that actual classes can implement -- in other words, these are not actual collections, just "prototypes" for the actual collections we will see in a moment.

The `Map` interface is not a subtype of `Collection` since maps work somewhat differently than other collections.

Concrete Collection Classes

- Behind the scenes, collections use well-known data structures to store elements

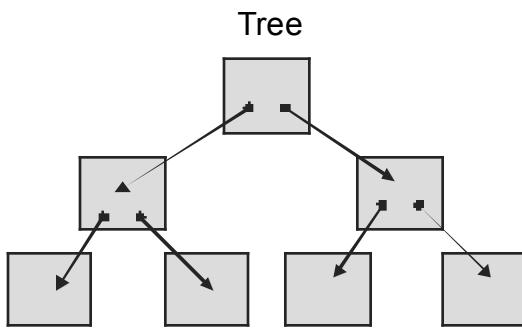
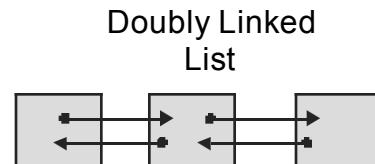
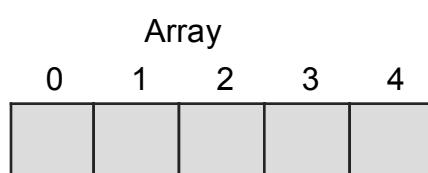


9 – 12

This figure shows a few of the "concrete" collection classes that Java provides; there are others. Note that each concrete class uses a different type of underlying data structure to implement the interface -- for example, an `ArrayList` uses an array, while a `LinkedList` uses a doubly-linked list.

Data Structures

- Each type of data structure has different performance, sorting and retrieval characteristics



Hash Table

Key	Value
Key	Value

9 – 13

Array-based collections have fast random access, but it's relatively slow to insert items in the middle, since the collection must move any existing elements back.

Linked lists have fast inserts, but relatively slow random access, since the collection must, on average, navigate through half of the collection to get to a given element (Java collections use doubly-linked lists).

Trees provide automatic sorting and relatively fast inserts and random access.

Hash tables use a key and a value -- you can look up values using the key relatively quickly.

Using an ArrayList

- **ArrayLists** provide fast random access of sequentially ordered elements

```
1  ArrayList<String> myList = new ArrayList<String>();  
2  
3  for ( int i = 0; i < 10; i++ )  
4      myList.add ( "adding element " + i );  
5  
6  for(String s : myList)  
7  {  
8      System.out.println(s);  
9 }
```

9 – 14

Here we show storing some String objects in an ArrayList and the so-called "enhanced for loop" to retrieve them.

Note: The Vector class works similarly to ArrayList -- the difference is that Vector is automatically thread-safe, while ArrayList is not.

Using an ArrayList, Pre Java 5

- Prior to Java 5, programs used **iterators** to visit a collection's elements

```
1  ArrayList myList = new ArrayList();
2
3  for ( int i = 0; i < 10; i++ )
4      myList.add ( "adding element " + i );
5
6  Iterator iter = myList.iterator();
7  while ( iter.hasNext() )
8  {
9      String s = (String)iter.next();
10     System.out.println ( s );
11 }
```

9 – 15

Java 5 adds "generic" collections, that lets you type a collection to the objects that will reside in the collection.

Using a LinkedList

- **LinkedLists** provide relatively fast inserts anywhere in the collection of sequentially ordered elements

```
1  LinkedList<String> linkList =  
2      new LinkedList<String>();  
3  
4  for ( int i = 0; i < 10; i++ )  
5      linkList.add ( "adding element " + i );  
6  
7  for(String s : linkList)  
8  {  
9      System.out.println(s);  
10 }
```

9 – 16

This code is almost the same as on the last page, but the performance is a little different. LinkedLists are faster than ArrayLists for inserts in the middle of the collection, but slower for random access.

So if your application needs fast random access and only adds elements at the end, ArrayList is a better choice than LinkedList. But if your program constantly needs to find a particular element, and then add a new element after it, LinkedList will be faster than ArrayList.

Using a HashSet

- **HashSets**, like all Sets, rejects duplicates

```
1  String[] beatles
2      = {"Paul", "George", "Ringo", "John", "Paul" };
3
4  HashSet<String> mySet = new HashSet<String>();
5  for ( int i = 0; i < beatles.length; i++ )
6  {
7      boolean added = mySet.add ( beatles[i] );
8      System.out.println ( "Added? " + added );
9  }
10
11 for(String s : mySet)
12 {
13     System.out.println(s);
14 }
```

9 – 17

Here we show attempting to add five strings to a HashSet, but since the last string is the same as the first, the last add will fail.

Note that to retrieve items from the HashSet, we can use the by-now familiar iterator.

Using a HashMap

- **HashMaps** provide keyed storage where the key values are unique

```
1  HashMap<String, String> myMap =  
2      new HashMap<String, String>();  
3  myMap.put ("John", "555-1111");  
4  myMap.put ("Paul", "555-2222");  
5  
6  String paulsNumber = myMap.get ("Paul");  
7  System.out.println (paulsNumber);  
8  
9  Set<String> keys = myMap.keySet();  
10  
11 for(String key : keys)  
12 {  
13     System.out.println(myMap.get(key));  
14 }
```

9 – 18

All Map-type collections store pairs of objects -- a key and a associated value. The utility of this is that you can thus associate a key with a value and quickly look up the value, given the key.

Like Sets, Maps reject duplicates; but with Maps, it's the key that must be unique.

Note that we can use a for loop to traverse through a Map, but we first must retrieve the collection of all of the keys.

Quick Practice

- Look in the JavaDocs for the `java.net.URL` class and note that it defines a constructor that takes a String url (e.g. "http://www.google.com"). Write a program fragment that creates an ArrayList and adds a few URL objects that reference your favorite Web sites. Then use a *for* loop that retrieves each URL object and prints it to the console.

Chapter Summary

In this chapter, you learned:

- Array and collection fundamentals
- How to code arrays
- The differences between various types of collections

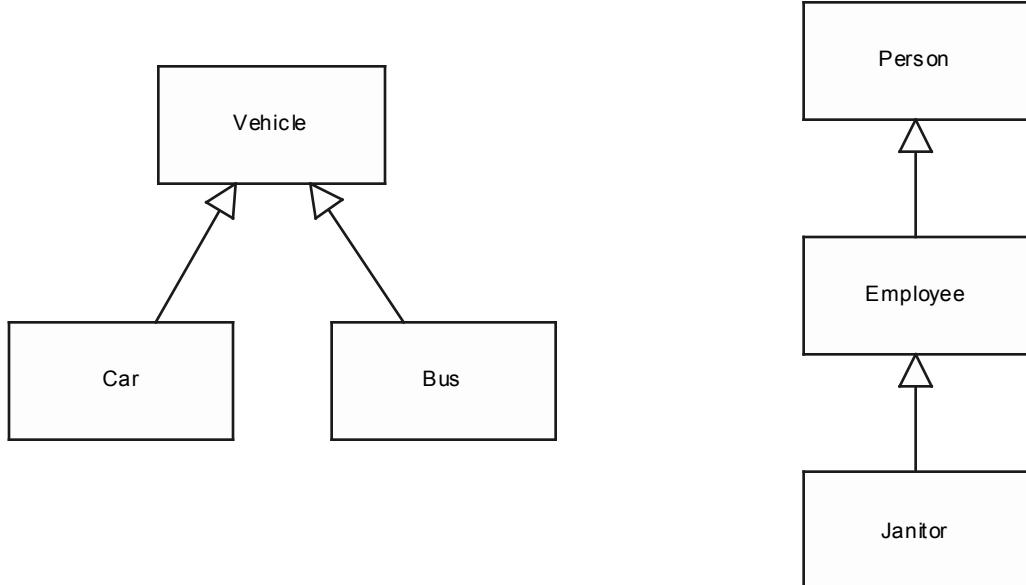
Inheritance and Polymorphism

- What is Inheritance?
- Inheritance versus Composition
- Polymorphism

10 - 1

What is Inheritance?

- In the real world, and in software systems, many things are "a kind of" another, more general thing



10 – 2

The notion of entities being a more specific thing than others of a similar type is common both in the real world and in software systems.

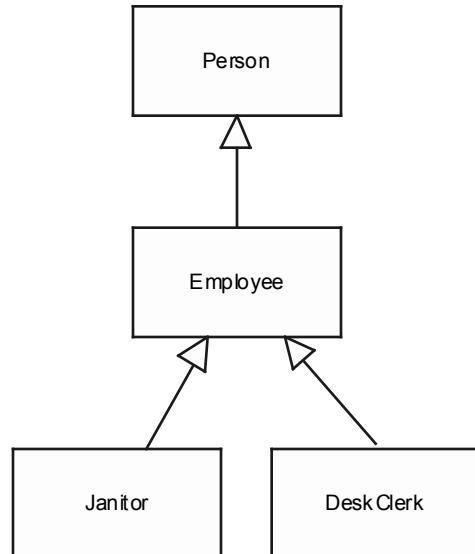
In a software system, we can model the "is-a" or "is-a-kind-of" relationship using inheritance. Inheritance is really just another kind of relationship in addition to the associations covered earlier.

The UML symbol for inheritance is a line with a broad-headed arrowhead. The "pointed-to" class is referred to as the superclass, while the other class is called the subclass.

Also remember that we model classes in UML, but what really counts in the relationships between objects of the classes.

Why Use Inheritance?

- If you can discern "is a kind of" relationships in your analysis, you can use inheritance to make your objects more reusable and extensible



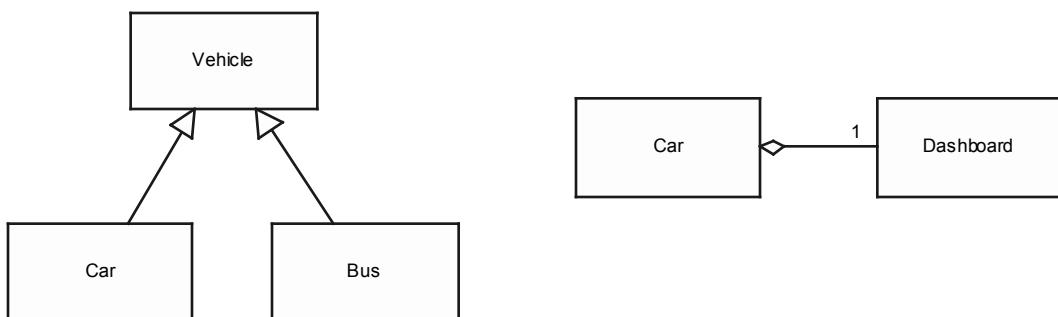
10 - 3

There are a couple of key benefits to using inheritance:

1. Inheritance lets you create models that closely approximate the real-world entity upon which the model is based. Lowering the so-called "representational gap" makes your software easier to understand and develop.
2. When combined with polymorphism (covered later), inheritance lets you create systems that are easy to extend without risking breakage to existing code.

Inheritance Vs Composed-Of Relationships

- We have seen classes that have aggregation or composition relationships with other classes
- You can generally determine this kind of relationship using the phrases "is a" (inheritance) versus "has a" (aggregation)

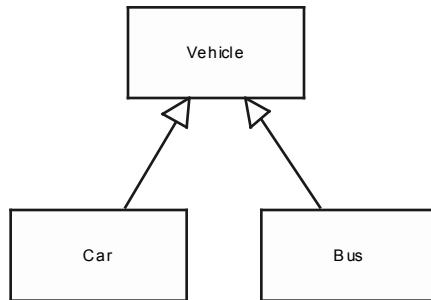


10 – 4

Sometimes it's difficult to determine whether to use inheritance or aggregation to implement a relationship, but you can use the "is-a" and "has-a" phrases as a quick test.

Java and Inheritance

```
1  public class Vehicle  
2  {  
3  }  
4  
5  public class Car extends Vehicle  
6  {  
7  }  
8  
9  public class Bus extends Vehicle  
10 {  
11 }
```

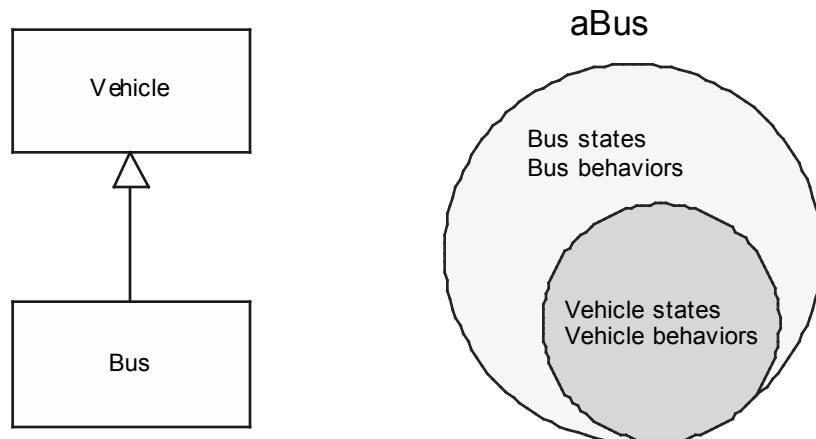


10 - 5

Java uses the keyword "extends" to implement an inheritance relationship.

Subclass Objects are a Superset

- Once instantiated, an object of a subclass type has all of the behaviors and states of its superclass in addition to any additional features defined by the subclass



10 – 6

This is the very basis of inheritance -- it allows you to create successively more refined and concrete classes that augment or replace states and behaviors from the superclass.

In this example, the additional state of a Bus might include the number of seats, whether the Bus has overhead storage, and so forth.

Accessing Superclass State and Behavior

- In most object-oriented programming languages, features marked as private are not accessible to subclasses
- That improves encapsulation and reduces the risk that changes in the superclass break subclasses

```
1  public class Vehicle
2  {
3      private int weight;
4  }
5
6  public class Bus extends Vehicle
7  {
8      public void myMethod()
9      {
10         System.out.println(weight); // ERROR!
11     }
12 }
```

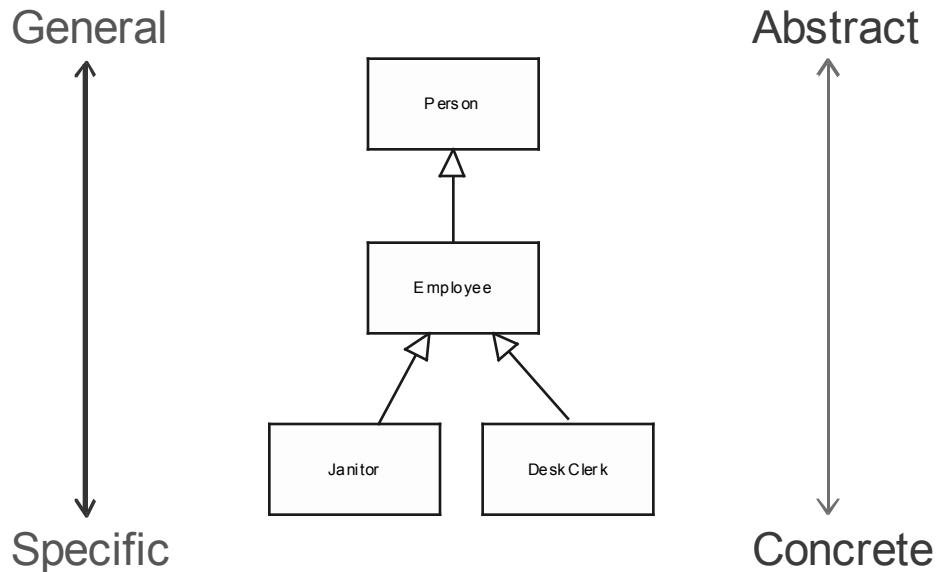
While it might seem weird that a subclass method cannot directly access states and behaviors defined in a superclass, the restriction actually makes for less brittle software. If a subclass can directly access features in a superclass, then the subclass might need to be changed if we change the feature in the superclass. That's unacceptable, especially considering that it's common for a different programmer to code the superclass and the subclasses.

So how would the Bus class access the Vehicle weight? The common practice is for the superclass to define public "get/set" methods that give controlled access to private members.

The downside of using "get/set" methods is that they are slower than direct access. So many object-oriented languages define an access level named "protected", that allows subclasses direct access. Protected access thus falls between private and public and provides weaker encapsulation, thus decreasing the malleability of the system, but perhaps increasing performance.

Building Class Hierarchies

- In a typical hierarchy, the superclasses are more general and less concrete than the subclasses



10 - 8

As you define subclasses, you add states and behaviors so that classes at the "bottom" of the hierarchy are more "real" than classes at the top.

As you design a hierarchy, you often "factor out" common states and behaviors and define them in a common superclass. This "factoring out" leads to designs that are easier to maintain, since if you need to change something, you only need to change it one place.

Quiz: What Kind of Relationship?

Shape

Ellipse

Polygon

Circle

Rectangle

Square

10 - 9

Quiz: What Kind of Relationship?

Book

Chapter

Paragraph

Sentence

10 - 10

Quiz: What Kind of Relationship?

George V

Edward VIII

George VI

Mary

John

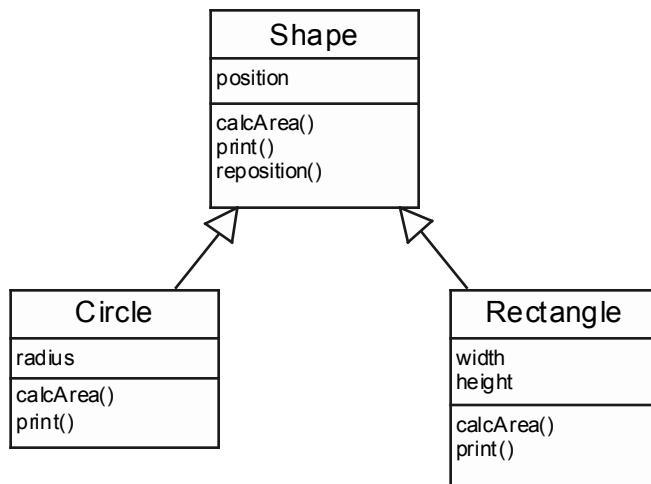
Elizabeth II

Margaret

10 - 11

Overriding Behaviors

- If a subclass provides a method with the same name and arguments as in the superclass, the subclass's method overrides the superclass method
- Overriding lets subclasses modify behaviors provided by superclasses



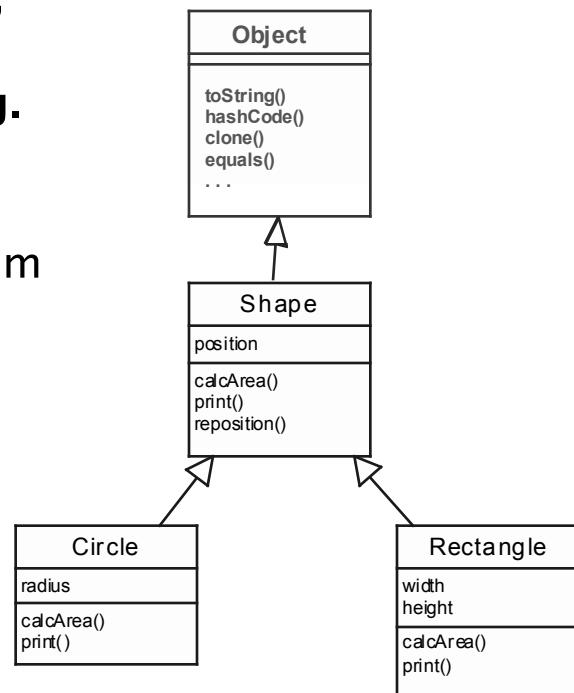
10 - 12

As we've already seen, a subclass can provide additional states and behaviors. A subclass can also CHANGE inherited behaviors by providing a method with the same name and arguments (method signature). When you create an instance of the subclass and call the method, the system invokes the method implementation defined in the subclass.

There are basically three things you can do in an overridden method: 1. Completely replace the implementation provided by the superclass(es). 2. Augment behavior provided by the superclass(es). In other words, do additional work, and also call back to the superclass implementation. 3. Completely remove the behavior by writing an empty method in the subclass.

The Object Class

- In Java, if you don't specify a superclass, Java uses the predefined **java.lang.Object** class
- All Java classes ultimately derive from Object



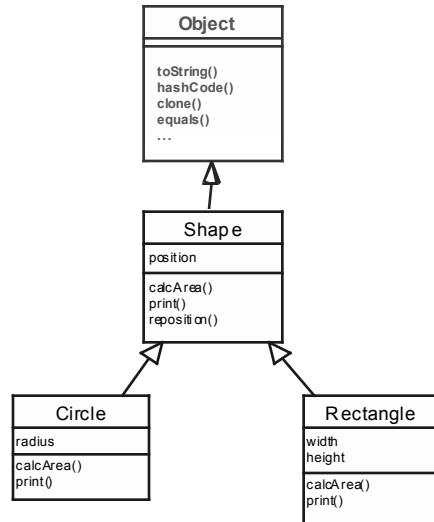
10 – 13

Object represents a generic Java object and provides the basic behavior required by all objects in Java.

The Object class is the only class that does not have superclass.

Calling Methods

```
1 int area = 0;
2 Point pos = new Point(2,3);
3
4 Shape s = new Shape();
5 Circle c = new Circle();
6 Rectangle r = new Rectangle();
7
8 s.reposition(pos);      // OK?
9 c.reposition(pos);      // OK?
10 r.reposition(pos);     // OK?
11
12 area = s.calcArea();   // OK?
13 area = c.calcArea();   // OK?
14 area = r.calcArea();   // OK?
15
16 String str = c.toString(); // OK?
```



10 - 14

Constructors and Superclasses

- Every constructor in a subclass must call a superclass constructor as its first line
- If you omit the call to the superclass, Java will insert it for you

```
1  public class Circle extends Shape
2  {
3      public Circle()
4      {
5          super();
6      }
7      public Circle(double radius)
8      {
9          System.out.println("1-arg ctor");
10     }
11     public Circle(Point position)
12     {
13         super(position);
14     }
15 }
```

Since constructors often have the job of initializing a class's fields, it's important that superclasses have the opportunity to initialize fields too. Therefore, Java requires that all constructors in subclasses, as their first executable line, invoke a constructor in the superclass using the `super()` keyword.

The subclass constructor can invoke ANY superclass constructor, not just a superclass constructor that has a signature matching the subclass constructor.

If you omit the call to `super()`, then Java automatically inserts a call to the superclass's no-argument constructor before any other executable statements in the subclass constructor.

If you explicitly call `super()`, the `super()` statement must be the first executable line or else the compiler will flag an error.

Constructors and Superclasses, cont'd

- Where's the compile error in this code?

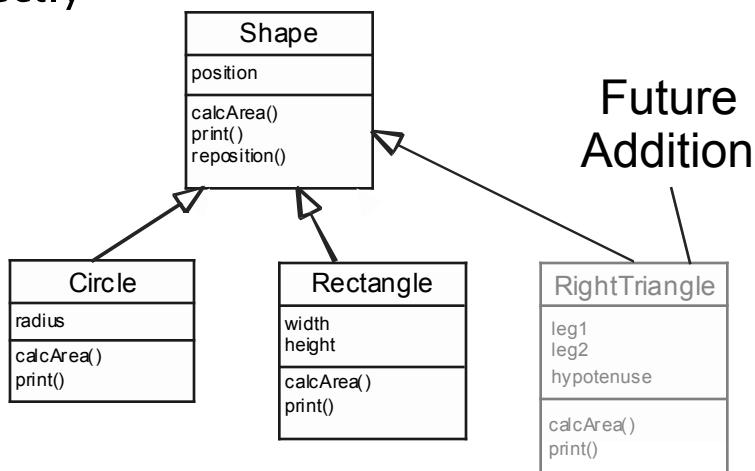
```
1  public class Shape
2  {
3      public Shape(Point position)
4      {
5          }
6      }
7
8  public class Rectangle extends Shape
9  {
10     public Rectangle()
11     {
12         System.out.println ("no-arg ctor");
13     }
14 }
```

10 - 16

This code will not compile since the Rectangle subclass's constructor doesn't have an explicit super() -- therefore the compiler assumes the no-argument constructor. Since the superclass (Shape) doesn't have a no-argument constructor, the compiler flags an error.

Polymorphism

- Polymorphism lets you design extensible, malleable systems to which you can add new classes without breaking existing code
- If you find yourself writing programs that use "if" statements to determine an object's type, you probably are not using polymorphism correctly



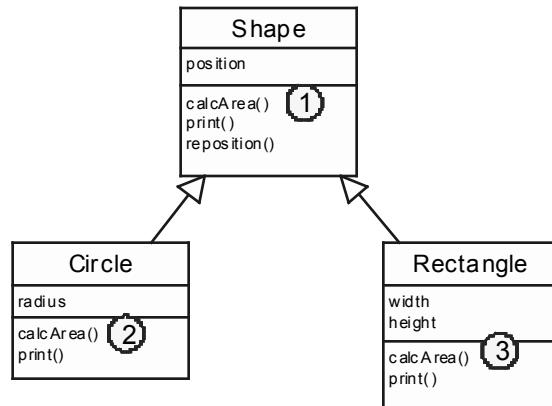
10 - 17

The idea behind polymorphism is that we often need to extend existing systems by adding new functionality and classes. Ideally, if we add a new class, we don't want the addition to require us to change any existing code.

The issue is that if our existing code uses some sort of "if" or "switch-case" statement to perform operations based on an object's type, then that code will need to be updated if we add a new class to the system. That's what polymorphism helps us avoid.

Polymorphic Reference Assignment

```
1  Shape p1 = new Shape();          // OK?
2  Circle p2 = new Circle();        // OK?
3  Rectangle p3 = new Rectangle();   // OK?
4
5  Shape p4 = new Circle();        // OK?
6  Shape p5 = new Rectangle();      // OK?
7  Circle p6 = new Shape();        // OK?
8
9  // which implementation runs?
10 p1.calcArea();
11 p2.calcArea();
12 p3.calcArea();
13 p4.calcArea();
14 p5.calcArea();
15 p4 = p5;
16 p4.calcArea();
```



10 - 18

Java allows you to assign references to subclass objects to variables typed as a superclass reference.

When you invoke a method on a reference, Java runs the subclass implementation. That means that at runtime, the JVM must examine the object's type, not the type of the reference. This is referred to as "late binding", since the chosen implementation is determined at runtime rather than at compile time.

Writing Polymorphic Algorithms

- Using polymorphism, you can write algorithms that work without modification even if you add new subclasses
- If you find yourself writing programs that use "if" statements to determine an object's type, you probably are not using polymorphism correctly

```
ArrayList<Shape> shapes =  
2     new ArrayList<Shape>();  
3     shapes.add(new Circle());  
4     shapes.add(new Rectangle());  
5  
6     double totalArea=0;  
7     for (Shape s : shapes)  
8     {  
9         totalArea += s.calcArea();  
10    }
```

10 - 19

The code shown here creates a Java ArrayList collection and then stores a few Circle and Square objects in the collection. Note that Circle and Square are subclasses of Shape.

Then to calculate the total area, we iterate (loop) through the collection, calling the calcArea() method on each Shape.

The key is that if we defined a new Shape, say a RightTriangle, the loop would still work the same! In other words, since all Shapes provide an overridden calcArea() method, our code doesn't break if we add a new kind of Shape. That's what polymorphism is all about.

Using Abstract Classes

- An abstract class cannot be instantiated, but can be subclassed
- Any class that contains any abstract methods must itself be marked as abstract

```
1  public abstract class Shape
2  {
3      private Point position;
4
5      public abstract double calcArea();
6      public void print()
7      {
8          System.out.println("Pos: " + position);
9      }
10 }
```

Shape s = new Shape();



10 – 20

An abstract method consists of a method definition with no body (i.e. no open and close curly braces). Any concrete subclass must provide an implementation of the method, or else the compiler will flag an error.

An abstract class can define data and can have non-abstract methods.

Classes with abstract methods let you define partially implemented classes high in a class hierarchy and force subclasses to provide required behaviors.

The chief differences between an abstract class and an interface is that interfaces can contain no code at all and that any given class can implement multiple interfaces, but can extend only one class.

The final Keyword

- A field marked as **final** is a constant
- A method marked as **final** in a superclass cannot be overridden in a subclass
- A class marked as **final** cannot be subclassed

```
1  public class A
2  {
3      public static final int MYVAL=12;
4
5      public final void myMethod()
6      {
7      }
8  }
9
10 public final class B
11 {
12 }
```

10 - 21

The author of a superclass can use the "final" keyword to prevent subclasses from overriding methods, or from even subclassing at all.

Quick Practice

- Write a class named **Policy** that represents an insurance policy. The class should have state for the policy number and amount. Then write a subclass named **AutoPolicy** that has additional state for the covered car model (String).
- If you have time, write constructors in each of the classes. The subclass constructor should invoke the superclass constructor.

10 – 22

Review Questions

- Why is good that subclass methods cannot access private state and behaviors from superclasses?
- What is the major benefit of polymorphism?

10 – 23

Chapter Summary

In this chapter, you learned:

- About the benefits of inheritance and class hierarchies
- The basics of using polymorphism

Interfaces

- What is an Interface?
- Implementing Interfaces

11 - 1

What is an Interface?

- Interfaces are a kind of **type** that help solve architectural issues:
 - Polymorphism without implementation inheritance (the Whistler example covered next)
 - As an architectural device so that an architect can specify a contract that programmers can implement
 - **Callbacks** from "system" software to application code
 - Remote computing using the **Proxy pattern**
 - And more...

11 – 2

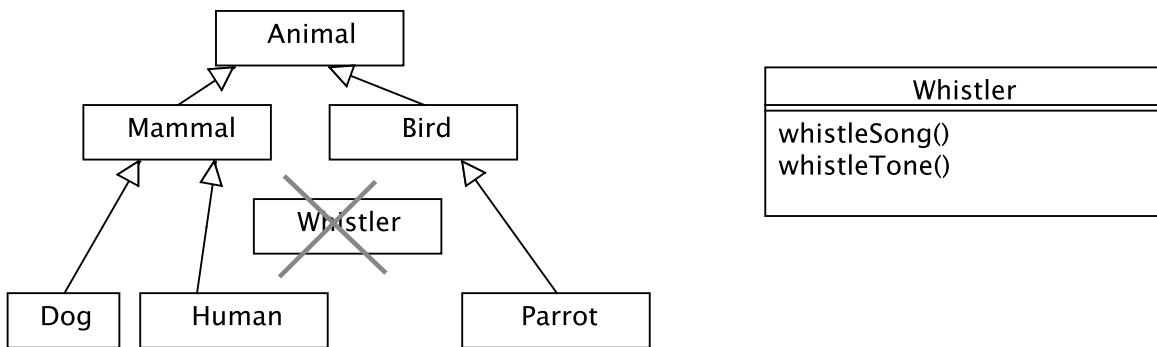
There is an article by Eben Hewitt available at:

<http://www.informit.com/articles/article.asp?p=331398&seqNum=2>

that describes other benefits.

Java Supports Single Inheritance

- Java only supports **single implementation inheritance** using the `extends` keyword
- Given the following diagram, how do we model **Whistlers**: both Humans and Parrots can whistle, but Dogs and other animals cannot



11 – 3

Some languages, notably C++, support multiple implementation inheritance, but Java does not. In other words, in Java, any given class can only have one superclass.

The designers of Java decided not to support it because it adds complexity and can lead to ambiguities in a class hierarchy. Look in the Wikipedia for "diamond inheritance problem" for an example of the issues.

We Still Need Polymorphism

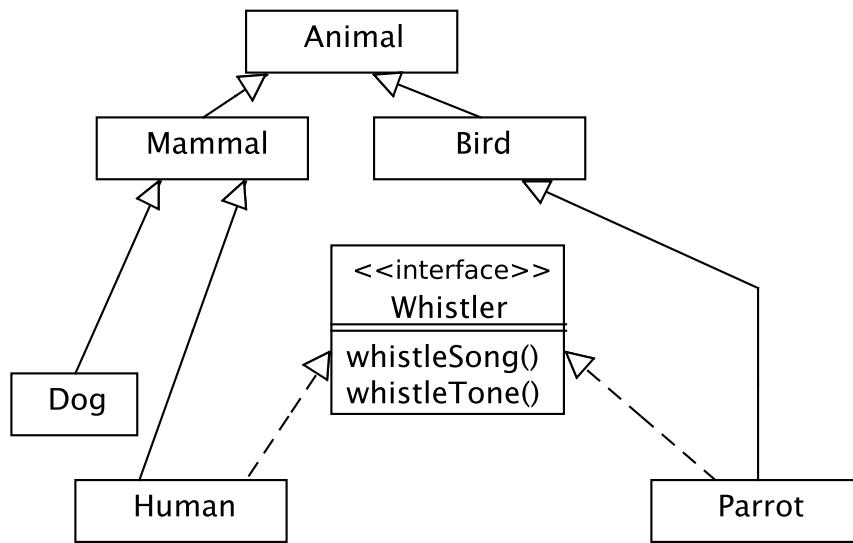
- Even though a Java class can only have one superclass, we would like to use polymorphism
- For example, we might want to create a list of **Whistlers** and traverse that collection

```
1  ArrayList<Whistler> whistlers =  
2      new ArrayList<Whistler>();  
3  whistlers.add(new Human());  
4  whistlers.add(new Parrot());  
5  . . .  
6  
7  for (Whistler w : whistlers)  
8      w.whistleTone(1000);
```

- How can we get this to work?

Solution: Interfaces

- A class can extend one superclass and **implement** zero or more **interfaces**



11 – 5

Interfaces effectively define a related set of behaviors that we refer to as a "role". A given class can play as many roles as necessary by provided the required behaviors (i.e. methods) that are defined in the interface.

One way to model interfaces in UML is to use the interface stereotype and then use a In the UML, one way to show a class implementing an interface is to draw a broad-headed arrow on a dashed line.

Software architects can define interfaces so that programmers can write classes that provide the required behavior.

Defining an Interface

- Java provides the **interface** keyword to define an interface
- An interface can define no method bodies, only method definitions

```
1  public interface Whistler
2  {
3      public abstract void whistleSong();
4      public abstract void whistleTone();
5  }
```

11 - 6

Interfaces contain zero or methods, each with no implementation. The methods can accept parameters, return a value, throw exceptions just like any Java method.

Interfaces can also define public fields as long as they are explicitly initialized in the field definition. Note that all interface fields are implicitly public, static and final.

Note that interface methods are automatically public and abstract, so it's not required to explicitly state that in the interface.

Implementing an Interface

- To compile, the implementing class must provide **all** of the interface methods

```
1  public class Human extends
2      Mammal implements Whistler
3  {
4      public void whistleSong() {...}
5      public void whistleTone() {...}
6      public void useTools() {...}
7  }
8
9  public class Parrot extends
10     Bird implements Whistler
11 {
12     public void whistleSong() {...}
13     public void whistleTone() {...}
14     public void fly() {...}
15 }
```

11 - 7

Since Human and Parrot implement Whistler, we can write the polymorphic code shown a few pages ago.

Classes that implement interfaces can also provide public methods that are not from the interface.

Implementing Multiple Interfaces

- A given class can implement any number of interfaces
- The implementing class must provide bodies for all of the interface methods

```
1  public class BurrowingOwl
2      implements Whistler, Digger
3  {
4      public void whistleSong() {...}
5      public void whistleTone() {...}
6      public void digHole() {...}
7      .
8  }
```

11 - 8

Although a given class can extend only one class, it can implement as many interfaces as required.

'Tagging' Interfaces

- The Java class library contains several interfaces that have no methods
- These so-called **tagging** interfaces act as a flag

```
public class ClassA implements java.io.Serializable
{
}
public class ClassB
{
}
1  public void testSerializable(Object o)
2  {
3      if (o instanceof java.io.Serializable)
4          serializeTheObject(o);
5      else
6          System.out.println ("Error: not serializable");
7 }
```

11 - 9

It's trivial to implement a tagging interface -- just write the "implements" clause on the class definition (there's no methods to implement!).

The `testSerializable()` method we wrote here lets us make a decision based on whether or not the specified object's class is tagged with the `Serializable` interface or not.

Interface Inheritance

- An interface can extend other interfaces
- Any class that implements a sub-interface must provide all of the methods from the sub and super interfaces

```
1  public interface InterfaceA
2  {
3      public void methodA();
4      public int methodB(double x);
5  }
6
7  public interface InterfaceB extends InterfaceA
8  {
9      public boolean methodC();
10 }
```

11 – 10

If an interface extends another interface, a class that implements the sub-interface must provide all of the methods defined in the specified interface and all of its super-interfaces.

Implementing a Sub-Interface

- To avoid compile errors, a class that **implements** an interface must provide at least empty implementations of all of the interface's **methods**

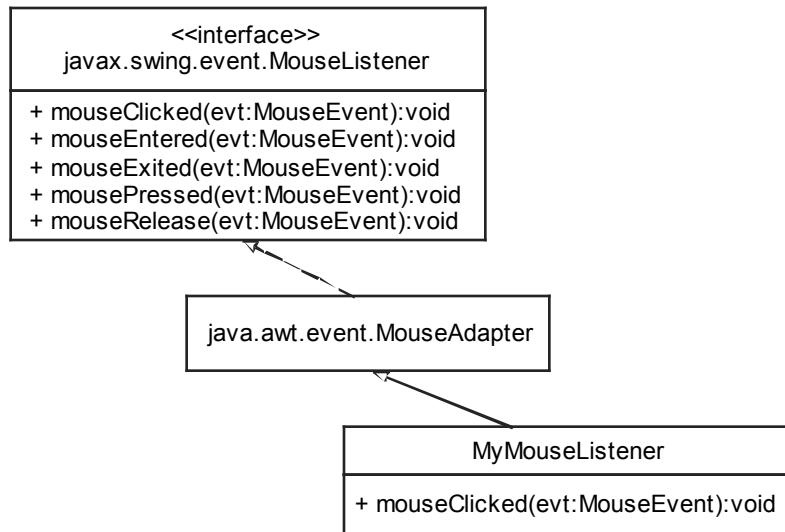
```
1 public class Test3 extends Test2
2     implements InterfaceB
3 {
4     public void methodA()
5     {
6     }
7     public int methodB(double x)
8     {
9         return 0;
10    }
11    public boolean methodC()
12    {
13        return false;
14    }
15 }
```

11 - 11

Java provides the "implements" keyword to show that a class provides the behavior specified by an interface and its super-interfaces (if any).

Convenience Adapter Classes

- In some cases, it's useful to define classes that implement interface methods with empty method bodies
- Then, subclasses can override only the methods that they desire



11 - 12

The Java class library provides several "adapter" classes that do nothing more than provide empty implementations for interface methods. That way, if you are writing a class that's only interested in a subset of the interface methods, you don't have to write a bunch of empty methods.

The downside is that a class can only extend one class, so if you extend an adapter, you cannot extend any other class. In contrast, a class can implement any number of interfaces.

Interfaces versus Abstract Classes

- An interface can have no method implementations or data
- A given class can implement multiple interfaces, but extend at most one class
- You cannot instantiate either an interface or an abstract class

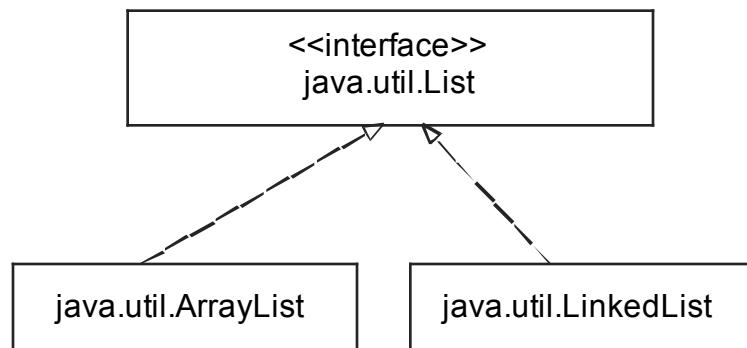
11 – 13

The key difference is that interfaces let designers specify sets of required behavior, while abstract classes (with abstract methods) let a class designer force a subclass to implement some or all of the methods.

Interface Reference Types

- Classes and interfaces are both considered to be **types**
- Java lets you assign subtype object references to supertype reference variables

```
List myList = new ArrayList();
```



11 - 14

Since you can say "is a", it's OK to assign a subtyped object reference to a variable that's typed to a super-type.

In this case, using the List supertype for the variable type is useful, since if we ever decide we'd rather use a LinkedList rather than an ArrayList, we only have to change the right side of the equals sign. In addition, we know that we will not get any compile errors since when we use a List-typed reference, we are restricted to using methods defined on the List interface (i.e., we are guaranteed that we did not use any methods specific to ArrayList).

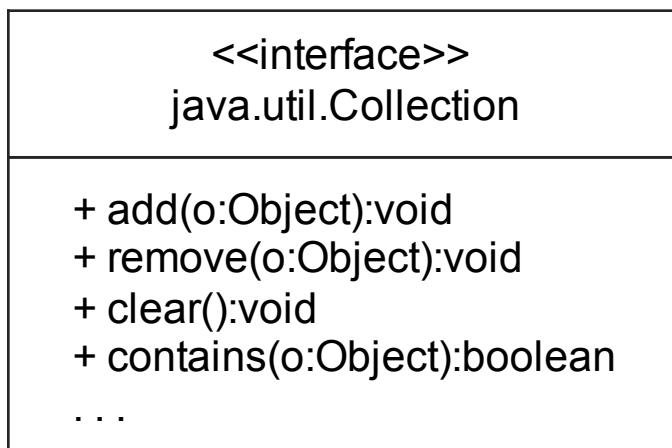
Quick Practice

- Look in the JavaDocs for the `java.awt.event.KeyListener` interface. Write a class that implements this interface. If you have time, add code to the class so that it echos the typed key value to the console.

11 - 15

Interfaces as Contracts

- Interfaces define a contract of required behavior (a role) that classes can implement
- For example, all Java collection classes ultimately implement the `java.util.Collection` interface, which defines the basic behavior of all collections



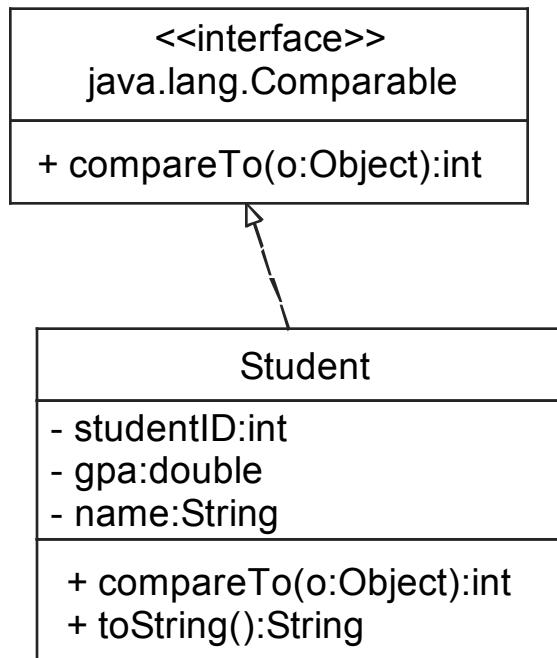
11 - 16

The Java Collections Framework defines a group of interfaces (e.g. `Collection`, `Set`, `List`) and provides a number concrete classes that implement those interfaces (e.g. `ArrayList`, `LinkedList` and `TreeSet`). The `Collection` interface is the super-interface that defines the behavior required of all Java collections.

Examples of Using Interfaces

- The rest of this chapter shows a couple of interesting uses of interfaces from the Java class library:
 - The **java.util.Comparable** interface
 - The **org.xml.sax.ContentHandler** interface

Example: Using the Comparable Interface



11 – 18

The Comparable interface from the `java.lang` package lets you indicate that objects in a class can play the role of comparing themselves to other objects. It's most often used to create sorted collections, as we will see in a moment.

Using the Comparable Interface, cont'd

```
1  public class Student implements Comparable
2  {
3      private int studentID;
4      private double gpa;
5      private String name;
6
7      public Student(int studentID,
8                      double gpa, String name)
9      {
10          this.studentID = studentID;
11          this.gpa = gpa;
12          this.name = name;
13      }
14      public String toString()
15      {
16          return "ID: " + studentID + " GPA: " + gpa +
17                  " Name: " + name;
18      }
```

11 - 19

Here we start the listing of the Student class. Note that it implements the Comparable interface.

Using the Comparable Interface, cont'd

```
19     public int compareTo(Object o)
20     {
21         Student s = (Student)o;
22         return this.studentID - s.studentID;
23     }
24 }
```

11 – 20

The Comparable interface defines a single method named "compareTo()", which is called on a given object, passing a reference to another object. The current object is supposed to compare itself to the other object.

This method returns an integer. According to the documentation for this interface, you should return a negative number if the current object is "less" than the other object, return zero if the objects are the same, and return a positive integer if the current object is "greater" than the other object.

When you write this method, you need to decide what it means to compare objects of the class. In this case, we decided that to compare Students, we compare their studentID values.

Using the Comparable Interface, cont'd

```
1 import java.util.Iterator;
2 import java.util.SortedSet;
3 import java.util.TreeSet;
4
5 public class Test
6 {
7     public static void main (String[] args)
8     {
9         SortedSet ss = new TreeSet();
10        ss.add (new Student(12, 3.44, "Sue"));
11        ss.add (new Student(8, 2.44, "Sam"));
12        ss.add (new Student(97, 1.44, "Cindy"));
13        ss.add (new Student(44, 4.44, "Bill"));
```

11 - 21

Now we show a class that uses the Comparable Student class defined on the last couple of pages.

Here we create a TreeSet, which is a concrete collection class that implements the SortedSet interface. As you might expect from its name, a SortedSet sorts items as they are added -- it determines where to place each new object in relation to existing objects by calling the compareTo method defined in the Comparable interface. (In fact, you will get an error if the added objects don't implement Comparable.)

Using the Comparable Interface, cont'd

```
15     Iterator iter = ss.iterator();
16     while (iter.hasNext())
17     {
18         Comparable c = (Comparable)iter.next();
19         System.out.println(c);
20     }
21 }
22 }
```

```
ID: 8 GPA: 2.44 Name: Sam
ID: 12 GPA: 3.44 Name: Sue
ID: 44 GPA: 4.44 Name: Bill
ID: 97 GPA: 1.44 Name: Cindy
```

11 - 22

Continuing the test program, we retrieve an iterator from the collection, and then simply print the objects. Note how we cast to Comparable!

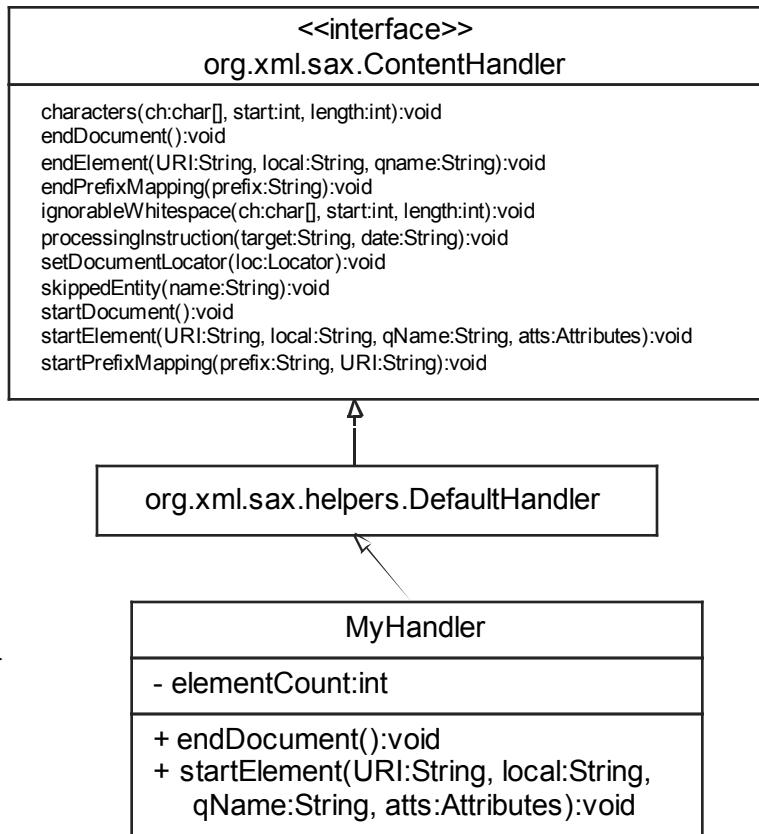
And as you'd expect, the output is sorted by studentID.

Example: Using the ContentHandler Interface

test.xml

```
<roster>
<student>
  <name>Bill Smith</name>
  <studentID>23</studentID>
  <gpa>3.44</gpa>
</student>
<student>
  <name>Sue Jones</name>
  <studentID>536</studentID>
  <gpa>3.99</gpa>
</student>
</roster>
```

11 - 23



Parsing XML content is an increasingly important part of Java programming. This example shows how to use the Simple API for XML Parsing (SAX) to count the number of elements in an XML file.

SAX is a "callback" API -- as the parser encounters XML content, the parser calls back to the program running the parser. For that callback to work, the handling class must define the methods expected by the parser. This is a archetypal use of interfaces: SAX defines the ContentHandler interface that lists the callback methods.

However, for simple parsing jobs, the ContentHandler interface provides too much information (but recall that if you implement an interface, you must provide ALL methods). So for a simple program like ours, we will instead subclass the DefaultHandler adapter class and only override the "startElement()" and "endDocument()" methods.

Example: Using ContentHandler, cont'd

```
1 import javax.xml.parsers.*;
2 import org.xml.sax.*;
3 import org.xml.sax.helpers.DefaultHandler;
4
5 public class MyHandler extends DefaultHandler
6 {
7     private int elementCount = 0;
8
9     public void endDocument() throws SAXException
10    {
11         System.out.println("Element count: " +
12                         elementCount);
13    }
14    public void startElement(String uri,
15                            String localName, String qName,
16                            Attributes attributes) throws SAXException
17    {
18        elementCount++;
19    }
11 - 24
```

Here we start the listing for the SAX-based parsing program that counts XML elements. The MyHandler class extends the DefaultHandler adapter class and overrides the startElement() and endDocument() methods. Note that since we were only interested in those two methods, it's more convenient to subclass DefaultHandler than to directly implement the ContentHandler interface, which has 11 methods.

The SAX parser calls startElement() each time it encounters an XML element -- we simply increment an element count integer. The SAX parser calls endDocument() only once -- when the parsing is complete. We use that opportunity to display the element count.

Example: Using ContentHandler, cont'd

```
20     public static void main (String[] args)
21     {
22         MyHandler handler = new MyHandler();
23         SAXParserFactory factory =
24             SAXParserFactory.newInstance();
25         try
26         {
27             SAXParser saxParser = factory.newSAXParser();
28             XMLReader xmlReader =
29                 saxParser.getXMLReader();
30             xmlReader.setContentHandler( handler );
31             xmlReader.parse( "test.xml" );
32         }
33         catch (Exception e)
34         {
35             e.printStackTrace();
36         }
37     }
38 }
```

11 - 25

Continuing the SAX element-counting program, here we show the main() method that initializes the parser and starts the parsing process.

The key line here is line 30, where we register an object of our MyHandler class as the callback object. The setContentHandler() method accepts a parameter of type ContentHandler, which works for us, since MyHandler extends DefaultHandler, which implements ContentHandler.

Review Questions

- What are the differences between an abstract class and an interface?
- Describe how a software architect might use interfaces while designing a software system

11 - 26

Chapter Summary

In this chapter, you learned:

- About the differences between interfaces and classes
- How to use interfaces to define roles that a class can implement